

---

# The Instrument Is the First Suspect: Six Broken Lenses in One Empirical Program

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Empirical claims about learned systems inherit the health of the instruments that  
2 measure them. We report six pre-publication incidents from one measurement-  
3 heavy research program on fast-weight sequence models in which a model, in one  
4 case a synthetic oracle, appeared to fail and the failure traced instead to the in-  
5 strument: an admission tolerance calibrated on the wrong key population; a probe  
6 reading a layer causal state-zeroing later proved inert; an identity-gauge assump-  
7 tion measured on one architecture and carried to another; an uncentered covari-  
8 ance estimator degenerate on the task’s near-orthogonal targets; a target whose  
9 ambient identity block taxed every rank budget; and a cross-check reference in-  
10 voked on the transposed side of a state convention. The same discipline adjudi-  
11 cated every incident, assembled across the catalogue and in full force before the  
12 final, decisive one: instrument-health adjudication before any model verdict, pre-  
13 registered crosschecks graded where the primary lens’s assumptions fail, falsifier  
14 controls run to completion, and raw-artifact tiebreaks. The discipline is itself falsi-  
15 fiable; its decisive crosscheck survived a pre-registered shuffled-target control that  
16 would have voided the adjudication. It does not launder results: the corrected lens  
17 left the endpoint verdict standing and two genuine model failures intact. We ar-  
18 gue the discipline transfers to any empirical evaluation gated by numerical checks,  
19 learned probes, or basis-dependent readouts.

## 20 1 Introduction

21 In one research program’s five-day window on fast-weight sequence models, six apparent headline  
22 model failures traced, under the same adjudication procedure, to broken instruments rather than  
23 models (Table 1); in two further readings it confirmed the model, not the lens, had failed. An  
24 incident here is a pre-registered endpoint or evaluation round read as a model verdict; defects caught  
25 outside such a reading (the auxiliary-classifier null of §4) are noted, not counted. An evaluation  
26 pipeline that never suspects its own instruments will publish lens artifacts as model results.

27 Our contributions: (1) the catalogue (Table 1), each incident with its mechanism and a reusable  
28 diagnostic signature; (2) a five-rule adjudication discipline (§2) that caught all six before publication;  
29 (3) the discipline’s boundary: it did not flip the pre-registered endpoint verdict, and two genuine  
30 model failures survived every instrument fix (§7).

## 31 2 The Adjudication Discipline

32 None of the five rules is novel in isolation; applied together, *before* any model verdict is recorded,  
33 they form a lens-failure detector with a measured record. Two clauses were codified mid-catalogue  
34 by the incidents that exposed them (Rule 3’s positive-control clause, §4; Rule 2’s graded crosscheck,  
35 Appendix A); all five rules were in force before the decisive adjudication of §5.

36 **Rule 1: instrument-health adjudication precedes model verdicts.** Before a sweep’s pre-  
 37 registered endpoint may be read, a health checklist walks the raw per-cell artifacts: gate violations,  
 38 convergence, a contradiction scan between independent readouts (the rule that fired in §5).

39 **Rule 2: pre-register a crosscheck that removes the primary lens’s strongest assumption,** graded  
 40 wherever that assumption provably fails, the switching rule fixed in advance.

41 **Rule 3: falsifier teeth, run to completion.** Every negative control is executed, not merely written;  
 42 guards are mutation-tested; every probe rung carries a planted-signal positive control; the adjudica-  
 43 tion itself carries a pre-registered falsifier whose firing voids it.

44 **Rule 4: contradictions are resolved against raw artifacts.** When two reads disagree, neither  
 45 recency nor authority decides: both are recomputed from the committed raw files and the tiebreak is  
 46 recorded before anything downstream runs.

47 **Rule 5: argmax closure.** Wherever a claim requires exact continuous recovery, the grading path  
 48 must not be satisfiable by argmax decoding, under which a rank-1 matrix can recover on the order  
 49 of  $d$  associations [Nichani et al., 2025]; the converse gap was pre-registered before it fired (§4).

### 50 3 Case I: The Tolerance That Manufactured a Cliff

51 A fast-weight associative memory’s recall capacity is measured by exact continuous recovery (cosine  
 52  $\geq 0.9$ , never argmax; Rule 5) across a load sweep  $K/d_{\text{state}}$ ; an eval-time admission gate decides  
 53 which trained cells enter a capacity fit; one leg checks that a Newton–Schulz orthogonalization in  
 54 the scoring path converges to residual  $< 0.01$  within  $n_{\text{iter}} = 20$  iterations. A wide-grid sweep at  
 55  $d_{\text{state}} = 96$  trained cleanly at every cell (training guard 1.0 throughout), yet only one of 12 new cells  
 56 was admissible ( $K=78, 84, 90$ : all three seeds failing): the shape of a capacity cliff, contradicting  
 57 the same architecture’s established super-linear growth at lower state sizes [Anonymous, 2026] and  
 58 the same grid’s own post-fix re-fit below.

59 The first mechanistic diagnosis blamed the trained anchor table’s convergence at high load; a follow-  
 60 up refuted it with margins (residuals near  $1.2 \times 10^{-6}$ ,  $7,800\times$  below tolerance, identical across  
 61 failing and passing cells; failures confined to a pool the anchor table bypasses) and the claim was  
 62 formally retracted.

63 A purpose-built repro instrument recovered the true failing tensor: the pool’s raw keys, never saved  
 64 by the original checkpoint writer. Its decision rule, a pre-registered strict precedence order, was  
 65 walked mechanically (Table 2): every instrument-bug branch failed to fire, and an iteration sweep  
 66 resolved all 4,608 flagged episodes by  $n_{\text{iter}} \leq 28$ . A wall keeps failing as the budget grows; a  
 67 miscalibrated tolerance resolves at a bump, and did. Root cause: a tolerance calibrated on anchor-  
 68 stabilized *blended* keys, silently applied to the structurally different *raw* keys this one pool queries.  
 69 Only the admission flag was recalibrated, only for the 11 cells verified to share the identical signa-  
 70 ture; the re-fit found no cliff through  $K/d_{\text{state}} = 0.9375$  (per-load mean recovery 0.92 to 1.00).

### 71 4 Case II: The Instrument at the Wrong Layer

72 A two-block fast-weight model is trained on an associative binding task ( $K = 32$  bindings per  
 73 episode) against matched baselines, with two registered readouts: a discrete task metric (episode-  
 74 restricted recall through the model’s own LM head) and a continuous probe metric  $\text{rf}@0.9$  (the  
 75 fraction of targets a linear probe on a tapped state recovers at cosine  $\geq 0.9$ ). In the third of three  
 76 consecutive rounds recorded as failures at  $\text{rf}@0.9 = 0.0$ , the discrete metric dissociated: recall  
 77 accuracy 0.9990 against chance 0.03125 (baselines 0.0447 and 0.0295), the probe still exactly 0.0  
 78 in every arm; the dissociation, not the zero, is the finding (Rule 5).

79 An offline re-fit battery removed the probe-training explanations: closed-form ridge regression,  
 80 pinned SGD, and MLP probes all reproduced  $\text{rf}@0.9 = 0.0$  exactly (over-regularization ruled  
 81 out by the rig’s post-nonlinearity recovery below); the trained probe’s output was the analytically  
 82 known episode-membership direction (mean cosine 0.176 against the  $1/\sqrt{32} \approx 0.177$  ceiling;  
 83 membership cosine 0.896); the ridge fit cleared its shuffled-tap control [Hewitt and Liang, 2019]  
 84 by at most +0.059 cosine. Meanwhile the full forward pass decodes the answer at 0.9957 accuracy  
 85 ( $31.9\times$  chance) for the contender, both baselines at or near chance: the answer exists in the network.




Figure 1: Primary versus crosscheck recovery (fraction recovered at cosine  $\geq 0.9$ , depth 64), all 62 cells. Non-converged cells (open orange) sit near the origin under both lenses; every cell with disagreement  $\geq 0.5$  is converged (filled blue). A lens defect that activates exactly where the model succeeds is the inverse of noise.

86 A state-zeroing experiment [Elazar et al., 2021] located the storage (Figure 2): zeroing the first  
 87 block’s state  $S_0$  collapses recall to chance (0.0286); zeroing  $S_1$ , the layer every deployed instrument  
 88 tapped, leaves it bit-unchanged at 0.9990: the probed layer is causally inert. Placement alone is not  
 89 the whole defect: a ridge probe on the load-bearing  $S_0$  *also* reads  $\text{rf}@0.9 = 0.0$  (shuffled-control  
 90 gaps  $+0.002$  to  $+0.063$ ), and only the post-block-1, pre-LM-head hidden exposes the recall  
 91 ( $\text{rf}@0.9 = 0.674$ , gap  $+0.800$ ), only in the arm that performs the task (ablation: 0.0; Table 3).  
 92 The bindings are stored in a format no tested linear read exposes; the model’s own downstream  
 93 nonlinearity is the only decoder found.

94 The same diagnosis exposed a second defect in passing: an auxiliary classifier whose labels were  
 95 uniform given the decoded quantity, so a perfect tap scores chance; its expected-null had been read  
 96 as confirmation for a round, hence Rule 3’s positive-control clause.

## 97 5 Case III: The Gauge That Did Not Transfer

98 A recurrent composer is trained on the word problems of five finite groups and scored for compo-  
 99 sitional depth generalization: trained at depth  $\leq 8$ , tested at depth 64. Two pre-registered lenses  
 100 score recovery of the group representation from the state: the *primary* de-gauges with a scale-only  
 101 correction, valid when the learned basis matches the target up to scale ( $\hat{Q} \approx I$ , measured on  
 102 every checkpoint of the *previous* architecture); the *crosscheck* fits a full orthogonal intertwiner by  
 103 Procrustes, fit and scored on disjoint splits, assuming nothing about the gauge.

104 **The signature.** The 62-cell sweep completed cleanly; run as built on the primary lens, the endpoint  
 105 machinery returned FALSIFY: no measurable separation from the baseline at the decisive groups.  
 106 The mandatory instrument-health pass (Rule 1) surfaced a contradiction (Figure 1): wherever  
 107 training had *not* converged, both lenses read near zero; wherever it converged, they contradicted  
 108 each other at full magnitude (one converged cell: loss 0.0001, primary 0.050, crosscheck 1.000).  
 109 All 13 cells with a lens disagreement of at least 0.5 are converged cells (threshold-robust: every cell  
 110 with a disagreement above 0.3 is converged, 16 of 16); on all 25 Arm-2 baseline cells (5 groups  $\times$   
 111 5 seeds, 50 checkpoint values) the lenses agree exactly. The harvest stopped for adjudication with  
 112 no model verdict recorded.

113 **The tiebreak.** A leaky crosscheck reading high on junk would leave the falsification standing. The  
 114 adjudication, against raw artifacts and code (Rule 4), rested on four grounds: (1) the crosscheck is  
 115 leakage-guarded by construction (fit and eval on disjoint splits); (2) it discriminates, reading near  
 116 zero on every non-converged cell; (3) the primary’s scale-only correction was already flagged as  
 117 basis-brittle across seeds (per-seed mean cosine 0.03 to 0.69 where the crosscheck read 0.86 to  
 118 0.95) and its  $\hat{Q} \approx I$  evidence came from a different architecture (a perfect model in a rotated basis  
 119 reads zero under an identity gauge); (4) the synthetic acceptance test already grades the crosscheck  
 120 when the injected gauge is non-identity (Rule 2). Verdict: the primary lens is broken on converged  
 121 cells; the mechanical falsification is void *as a model verdict*.

122 **Teeth.** The tiebreak carried its own pre-registered falsifier (Rule 3): shuffled-target re-scoring must  
 123 read near zero, or the tiebreak is wrong and the contradiction escalates to an instrument rebuild.  
 124 On three pre-registered converged checkpoints spanning three groups (among them the converged  
 125 exemplar above), the shuffled crosscheck read 0.00, 0.00, and 0.05, far below the 0.5 void threshold,  
 126 while the unshuffled reads reproduced the committed values bit-identically (Table 5).

127 **The re-metric.** The full grid was re-read under the crosscheck lens through the unmodified verdict  
 128 machinery, which first reproduced the committed primary verdict and all 62 recomputed ceilings  
 129 to within  $10^{-9}$ . The corrected endpoint keeps the same FALSIFY verdict *and the same machine-*

130 *emitted reason* (§7, Table 4); what the fix changes is the readable geometry: the primary was not  
 131 a blanket null (S3 cleared its bar, 0.54 against 0.459) but masked which groups fail, while under  
 132 the crosscheck the two decisive failures, an S5 near-miss and an A6 non-convergence, stand out  
 133 lens-independently. Fixing the instrument changed what the verdict means, not the verdict.

## 134 6 Three More Lenses, Briefly

135 Three further incidents (rows 4–6 of Table 1; numbers in Appendix A): an **uncentered covariance**,  
 136 degenerate on the task’s own near-orthogonal targets, that failed a synthetic *perfect* model (the  
 137 rescue de-confounds into centering plus a full-intertwiner degauge); a **target with an ambient**  
 138 **identity block**, under which a causal rank experiment returned near-chance recovery at every rank  
 139 (zero in four of five groups; 0.17 and 0.08 in the fifth) while the capped arms sat at the exact rank- $k$   
 140 ceiling  $\sqrt{k/d_{\text{state}}}$ : undelivered by the instrument, not failed by the model; and a **reference kernel**  
 141 **invoked on the transposed side** of a state convention, “contradicting” a correct recurrence at the  
 142  $\sqrt{2}$  signature of a transposed outer product, exonerated by a closed-form hand computation.

## 143 7 The Boundary: What the Discipline Did Not Flip

144 A lens-repair procedure that only rescued the preferred hypothesis would be a laundering machine;  
 145 under the corrected lens, Case III’s endpoint remained **FALSIFY**, held there by two lens-independent  
 146 model results: one decisive configuration never reaches training-support convergence (recovery 0.00  
 147 at depth 64, all five seeds, both lenses), and the other decisive group sits below its pre-registered bar  
 148 (0.483 against 0.735) lens-independently: one seed generalizes at 0.00 under both lenses despite low  
 149 training loss, and even its siblings reach only 80% and 65% of their own ceilings, under the 90%  
 150 criterion, so the group misses its bar with or without the outlier, against a baseline pinned at 0.0.

151 Case I shows the same refusal: a diagnosed pool-margin confound ( $+0.0330$ ,  $12\times$  the 0.0028 noise  
 152 floor) was withheld because the diagnostic’s own cells never cleared admission, and an archived  
 153 exact-ceiling reading (1.0000, three seeds) that did not replicate at a fresh seed (0.9725) was re-  
 154 scoped to the seeds measured; the machinery that unlocked eleven quarantined cells declined, twice,  
 155 to promote findings it could not support.

## 156 8 Related Work and Limitations

157 **Related work.** Silberzahn et al. [2018] and D’Amour et al. [2022] document divergent conclusions  
 158 from underspecified pipelines; here the models are fixed and the *measurement pipeline* is the adjudi-  
 159 cated object. Henderson et al. [2018] and Bouthillier et al. [2021] quantify seed and tuning variance;  
 160 none of the six defects is sampling variance, so variance accounting catches none. REFORMS  
 161 [Kapoor et al., 2023], Sculley et al. [2018], and Lipton and Steinhardt [2018] prescribe what a  
 162 trustworthy paper discloses; we supply the in-flight machinery such checklists presume exists.  
 163 The probing literature supplies the instrument-hygiene checks the wrong-layer case rediscovers at  
 164 evaluation scale: control tasks and selectivity [Hewitt and Liang, 2019], the probe-versus-behavior  
 165 dissociation [Belinkov, 2022], amnesic probing [Elazar et al., 2021], and randomization-based  
 166 sanity checks [Adebayo et al., 2018]. Schaeffer et al. [2023] argue the inverse of our thesis (an  
 167 apparent model *property*, emergence, as metric artifact) from the same instrument-first stance; each  
 168 diagnoses one instrument class, where this catalogue is a serial, pre-registered, cross-instrument  
 169 adjudication record with falsifier teeth. Nichani et al. [2025] prove Rule 5’s argmax/exact-recovery  
 170 gap; we report it firing live (§4); the fast-weight testbeds follow Schlag et al. [2021].

171 **Limitations.** All six incidents come from one program, one team, one architecture family: they  
 172 support no prevalence claim, and we cannot report the rate a lighter process would miss. The  
 173 catalogue is survivor-biased: a defect missed by both crosscheck and falsifier battery appears in no  
 174 record we could write. It is biased by trigger too: the discipline fires on apparent model failure, so  
 175 defects flattering the model are under-represented; five of six fixes moved results in the program-  
 176 favorable direction. In the one incident costed, the two diagnostic instruments took 0.50 and 1.10  
 177 GPU-hours (archived wall-clock) against the 6.33 GPU-hours of training the uncorrected readings  
 178 would have discredited. We report no false-alarm rate: the records do not tally escalations against  
 179 clean instrument-health passes, so the precision of “suspect the instrument first” is unmeasured.

## References

- Julius Adebayo, Justin Gilmer, Michael Muelly, Ian Goodfellow, Moritz Hardt, and Been Kim. San-  
ity checks for saliency maps. In *Advances in Neural Information Processing Systems (NeurIPS)*,  
2018. arXiv:1810.03292.
- Anonymous. The capacity cliff is not capacity. Companion paper, under review; citation anonymized  
for double-blind review. Reports the substantive capacity results whose measurement incidents  
this paper analyzes., 2026.
- Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational  
Linguistics*, 48(1):207–219, 2022. arXiv:2102.12452.
- Xavier Bouthillier, Pierre Delaunay, Mirko Bronzi, Assya Trofimov, Brennan Nichyporuk,  
Justin Szeto, Nazanin Mohammadi Sepahvand, Edward Raff, Kanika Madan, Vikram Voleti,  
Samira Ebrahimi Kahou, Vincent Michalski, Tal Arbel, Chris Pal, Gaël Varoquaux, and Pascal  
Vincent. Accounting for variance in machine learning benchmarks. In *Proceedings of Machine  
Learning and Systems (MLSys)*, 2021. arXiv:2103.03098.
- Alexander D’Amour, Katherine Heller, Dan Moldovan, Ben Adlam, Babak Alipanahi, Alex Beutel,  
Christina Chen, Jonathan Deaton, Jacob Eisenstein, Matthew D. Hoffman, et al. Underspecifica-  
tion presents challenges for credibility in modern machine learning. *Journal of Machine Learning  
Research*, 23(226):1–61, 2022. arXiv:2011.03395.
- Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. Amnesic probing: Behavioral  
explanation with amnesic counterfactuals. *Transactions of the Association for Computational  
Linguistics*, 9:160–175, 2021. arXiv:2006.00995.
- Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger.  
Deep reinforcement learning that matters. *Proceedings of the AAAI Conference on Artificial  
Intelligence*, 32(1), 2018. arXiv:1709.06560.
- John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Empirical  
Methods in Natural Language Processing (EMNLP)*, 2019. arXiv:1909.03368.
- Sayash Kapoor, Emily Cantrell, Kenny Peng, Thanh Hien Pham, Christopher A. Bail, Odd Erik  
Gundersen, Jake M. Hofman, Jessica Hullman, Michael A. Lones, Momin M. Malik, Priyanka  
Nanayakkara, Russell A. Poldrack, Inioluwa Deborah Raji, Michael Roberts, Matthew J. Sal-  
ganik, Marta Serra-Garcia, Brandon M. Stewart, Gilles Vandewiele, and Arvind Narayanan.  
REFORMS: Reporting standards for machine learning based science. *arXiv preprint  
arXiv:2308.07832*, 2023.
- Zachary C. Lipton and Jacob Steinhardt. Troubling trends in machine learning scholarship. In *ICML  
2018 Machine Learning Debates Workshop*, 2018. arXiv:1807.03341.
- Eshaan Nichani, Jason D. Lee, and Alberto Bietti. Understanding factual recall in transformers via  
associative memories. In *International Conference on Learning Representations (ICLR)*, 2025.  
arXiv:2412.06538.
- Rylan Schaeffer, Brando Miranda, and Sanmi Koyejo. Are emergent abilities of large language  
models a mirage? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.  
arXiv:2304.15004.
- Imanol Schlag, Kazuki Irie, and Jürgen Schmidhuber. Linear transformers are secretly fast  
weight programmers. In *International Conference on Machine Learning (ICML)*, 2021.  
arXiv:2102.11174.
- D. Sculley, Jasper Snoek, Alex Wiltschko, and Ali Rahimi. Winner’s curse? on pace, progress,  
and empirical rigor. In *International Conference on Learning Representations (ICLR) Workshop  
Track*, 2018. OpenReview: rJWF0Fywf.
- Raphael Silberzahn, Eric Luis Uhlmann, Daniel P. Martin, Pasquale Anselmi, Frederik Aust, Eli  
Awtrey, Štěpán Bahník, Feng Bai, Colin Bannard, Evelina Bonnier, et al. Many analysts, one data  
set: Making transparent how variations in analytic choices affect results. *Advances in Methods  
and Practices in Psychological Science*, 1(3):337–356, 2018.

## A The Three Brief Lenses, in Full

**The covariance that could not see.** An acceptance test for a representation-recovery pipeline injected a synthetic *perfect* model and failed it. The subspace step estimated an uncentered covariance, degenerate on the task’s own targets: for near-orthogonal states  $Z \approx c(\rho \oplus I)$ ,  $ZZ^\top \approx c^2 I$  is isotropic and carries no subspace information. On the production harness, the headline rescue of the oracle injection (mean cosine 0.084 to 0.9996) folds in two registered lens changes: centering *and* a switch from the scale-only primary to the full-intertwiner crosscheck. Isolating centering under a matched full-intertwiner metric moves mean cosine 0.705 to 0.9996; under the scale-only lens the centered pipeline reads 0.046, no better than the uncentered 0.084, because the injection’s true gauge is non-identity, the pre-echo of §5. Centering is a genuine fix, but the full rescue is centering plus a full-intertwiner degauge, not centering alone. The fix (centering) cancels the constant identity block because the group mean of a nontrivial representation vanishes; the acceptance injection was extended so the previously skipped subspace step now sits inside the test.

**The target that taxed every budget.** A causal test of a rank-recruitment law capped model rank at  $k$  around the predicted minimum  $d_{\min}$  and expected recovery to return at  $k = d_{\min}$ . Every capped arm returned near-chance recovery on both lenses: zero in four of the five groups, and 0.17/0.08 in the fifth (S3), far below that group’s ceiling. The target had been constructed as  $\rho \oplus I$ : the ambient identity block inflates the as-built target rank by two, so every rank budget buys the constant block first. The diagnosis is quantitative: a rank- $k$  model’s best achievable cosine against this target is exactly  $\sqrt{k/d_{\text{state}}}$ , and 37 of 39 force-rank cells sat within 0.07 of that ceiling (mean absolute deviation 0.028); the arms had trained to their theoretical optimum against the wrong target. The causal test was undelivered by the instrument, not failed by the model; a corrected-target, zero-padded re-run subsequently delivered the test (companion work, anonymized) [Anonymous, 2026].

**The reference on the wrong side.** A bespoke recurrence was cross-checked against a reference kernel and disagreed at relative Frobenius residual 1.40; the recurrence was the initial suspect. At depth one the update is closed-form ( $S_1 = \beta v k^\top$ ), so a hand computation with scalar loops adjudicates conventions with no numerics: the composer matched the closed form at  $4.5 \times 10^{-8}$  and its *transpose* at 1.405, and for near-orthogonal unit  $k$  the transpose signature is exactly  $\|v k^\top - k v^\top\|_F / \|v k^\top\|_F \approx \sqrt{2}$ . The reference library’s documented state convention is the transpose of the pinned equation; the cross-check’s own reference invocation was the defect. With the returned state transposed before comparison, the cross-check passes all three pinned configurations (residuals  $2.8 \times 10^{-3}$  to  $4.5 \times 10^{-3}$ ), and a permanent negative test confirms a transposed-update mutant is killed at 1.405.

## B The Catalogue and Per-Incident Detail Tables

Table 1 is the six-incident catalogue in one view; the remaining tables and Figure 2 carry the per-incident numbers the body summarizes. Every value traces to the same evidence rows (and raw artifacts, with checksums) recorded in the paper’s claims-to-evidence map.

| Apparent model failure                           | Actual instrument defect                                              | Diagnostic signature                                                | Resolution                                               |
|--------------------------------------------------|-----------------------------------------------------------------------|---------------------------------------------------------------------|----------------------------------------------------------|
| Capacity collapse, 11/12 cells quarantined (§3)  | tolerance calibrated on blended keys, applied to raw keys             | every flagged episode resolves under an iteration bump              | flag recalibrated; no cliff                              |
| Three failed evaluation rounds (§4)              | probe tap on a causally inert layer; storage not linearly decodable   | task metric near ceiling, probe zero; state-zeroing localization    | metric re-registered; tap re-located                     |
| Generalization endpoint falsified (§5)           | identity gauge from one architecture carried to another               | primary/crosscheck contradiction, only on converged cells           | crosscheck decisional; verdict re-read (still falsified) |
| Perfect synthetic model fails acceptance (§6)    | uncentered covariance, degenerate on near-orthogonal targets          | oracle injection scores near zero; centering restores it            | estimator centered; injection extended                   |
| Causal rank test near-chance at every rank (§6)  | target built with an ambient identity block, taxing every rank budget | capped cells sit at the exact ceiling $\sqrt{k/d_{\text{state}}}$   | target zero-padded; test re-run                          |
| Recurrence contradicted by reference kernel (§6) | reference invoked on the transposed side of a state convention        | relative residual $\approx \sqrt{2}$ ; closed-form hand computation | reference transposed; mutant-kill check added            |

Table 1: The six incidents: in each, the model was initially the suspect and the instrument was the finding, and every row was adjudicated against raw artifacts before any model-level conclusion was recorded. Sections named per row.




Figure 2: Case II, the wrong-layer instrument adjudicated causally. (a) Episode-restricted recall accuracy under state zeroing, contender (blue) and ablation baseline (gray), chance = 1/32 dashed: zeroing  $S_0$  collapses recall to chance in both arms; zeroing  $S_1$ , the layer every deployed probe tapped, changes nothing. (b) Linear (ridge) probe recovery rf@0.9 by tap site: both recurrent states read zero in both arms, including the causally load-bearing  $S_0$ ; only the post-nonlinearity pre-LM-head hidden exposes the recall, and only in the arm that can perform the task. The figure regenerates from the archived localization summary under an asserted checksum.

| Step (pre-registered order)    | Outcome  | Reading                                                   | Rules out                   |
|--------------------------------|----------|-----------------------------------------------------------|-----------------------------|
| entity-set reconstruction gate | pass     | 107/107 set-equal                                         | wrong-partition analysis    |
| pool-membership precheck       | pass     | 0 violations / 36 events                                  | instrument bug (structural) |
| reproduction-count floor       | pass     | 12× the minimum                                           | fluke event                 |
| live/offline agreement         | pass     | 0 disagreements                                           | instrument bug (numerical)  |
| iteration sweep                | resolves | 4,313 @ $n_{\text{iter}}=24$ ; 295 @ $n_{\text{iter}}=28$ | structural wall             |

Table 2: Case I, the repro instrument’s decision-rule walk on the real cell ( $K=84$ ): each step ran in its pre-registered precedence order, and the terminal verdict (tolerance miscalibration) was reachable only after every instrument-bug branch failed to fire. All 4,608 flagged episodes resolve by  $n_{\text{iter}} = 28$ ; none remained at larger budgets. Evidence row I3.

| tap site                             | contender    |               | ablation |        |
|--------------------------------------|--------------|---------------|----------|--------|
|                                      | rf@0.9       | gap           | rf@0.9   | gap    |
| (i) $S_1$ , shallow query (deployed) | 0.0          | +0.060        | 0.0      | +0.005 |
| (ii) $S_0$ , bind pathway            | 0.0          | +0.006        | 0.0      | +0.002 |
| (iii) $S_1$ , true query pathway     | 0.0          | +0.063        | 0.0      | +0.005 |
| (iv) pre-LM-head hidden              | <b>0.674</b> | <b>+0.800</b> | 0.0      | +0.006 |

Table 3: Case II, ridge-probe recovery by tap site; “gap” is each fit’s cosine-mean margin over its own shuffled-tap control. No state-level linear tap recovers anything in either arm, including (ii) on the causally load-bearing state; the post-nonlinearity tap (iv) recovers strongly and only in the arm that performs the task. Evidence row W4.

| group         | ceiling (X) | far (X) | bar   | clears | baseline far (X) | separates |
|---------------|-------------|---------|-------|--------|------------------|-----------|
| S3            | 0.920       | 0.900   | 0.828 | yes    | 0.150            | no        |
| S4            | 1.000       | 1.000   | 0.900 | yes    | 0.000            | yes       |
| A5            | 0.490       | 0.390   | 0.441 | no     | 0.000            | no        |
| S5 (decisive) | 0.817       | 0.483   | 0.735 | no     | 0.000            | no        |
| A6 (decisive) | 0.190       | 0.000   | 0.171 | no     | 0.000            | no        |

Table 4: Case III, the re-metric: the pre-registered endpoint computed under the crosscheck lens (X) per group, by the unmodified verdict machinery on shadow cells; “far” is depth 64 and “bar” is 90% of each group’s own ceiling. The endpoint remains FALSIFY (neither decisive group separates), but the failure geometry is now readable: A6 never converges at its decisive configuration under either lens, and S5 misses its bar lens-independently (one seed generalizes at 0.00; its two siblings reach only 80% and 65% of their own ceilings, both below the 90% bar). Separation requires the contender to clear its bar *and* the Arm-2 baseline to collapse below 50% of its own ceiling; S3 clears its bar but its baseline does not collapse ( $0.15 > 0.125$ ), hence no separation. Evidence row X5.

| converged control checkpoint | real                  | shuffled | falsifier ( $\geq 0.5$ ) fires |
|------------------------------|-----------------------|----------|--------------------------------|
| A6, $n_h=4$ , seed 0         | 1.000 (bit-identical) | 0.000    | no                             |
| S5, $n_h=4$ , seed 0         | 0.800 (bit-identical) | 0.000    | no                             |
| S4, $n_h=2$ , seed 2         | 1.000 (bit-identical) | 0.050    | no                             |

Table 5: Case III, the tiebreak’s own pre-registered falsifier: on three converged checkpoints, the crosscheck re-scored with targets shuffled across the item pool reads far below the 0.5 void threshold, while the unshuffled reads reproduce the committed values exactly (“bit-identical”). Had any shuffled read cleared 0.5, the tiebreak was pre-registered to be wrong and the adjudication void. Evidence row X4.