

Constant-Memory Recall: A Fixed 32 kB Fast-Weight State Retains In-Context Bindings That a Matched Transformer Does Not Learn

Anonymous authors

Paper under double-blind review

Abstract

Long-horizon inference is increasingly memory-bound: a transformer’s key-value (KV) cache grows linearly with context length, while fast-weight recurrences carry a state of fixed size. Whether such a state retains the bindings a recall task depends on, under a training budget an attention model also receives, is an empirical question. We compare three models matched in parameters, data, and training schedule on episode-restricted 32-way associative recall: a two-block matrix fast-weight model holding 32,768 bytes of recurrent state, a flat-vector recurrence, and a transformer. The fast-weight model answers recall queries at 0.9995 mean accuracy over three seeds; both baselines sit at chance, and paired-seed confidence intervals exclude the pre-registered 0.30 separation margin with lower bounds above 0.95. Accuracy stays at or above 0.998 for every seed when the same episode is embedded in contexts of 454 to 1,798 tokens, eight times the binding span, with the state fixed at 32,768 bytes; the transformer reads chance uncapped and at every KV-cache budget from one to thirty-two times those bytes, so no memory multiplier is claimed and the baseline is reported as non-competitive at the matched budget. State-zeroing localizes the recall to the first block’s fast-weight state. A multi-hop variant separates only as seed variance, and at higher load every arm reads chance.

1 Introduction

Serving a transformer over a long context is a memory problem before it is a compute problem. The KV cache of a decoder-only model (Vaswani et al., 2017) grows linearly with context length, and at deployment scale that growth, not arithmetic throughput, sets the limit on batch size, context window, and cost. A substantial literature therefore compresses or evicts cache entries (Xiao et al., 2023; Zhang et al., 2023), trading recall fidelity for bytes held.

Fast-weight recurrences occupy the opposite corner of this trade. A model whose mixer maintains a matrix state updated by an outer-product rule (Schmidhuber, 1992; Ba et al., 2016; Schlag et al., 2021; Katharopoulos et al., 2020; Yang et al., 2024b; Gu & Dao, 2023) carries a memory of fixed byte size regardless of context length. The architecture promises constant-memory inference, but the promise rests on an empirical premise that is rarely tested head to head: that the fixed state, trained under the same budget an attention model receives, in fact retains the bindings a long-horizon task queries. Prior diagnostic work has, if anything, pointed the other way; on multi-query associative recall, attention models dominate efficient recurrent alternatives when both are trained well (Arora et al., 2023).

We report a matched-budget comparison in which the fixed state wins, and wins at ceiling. On episode-restricted 32-way associative recall, a two-block matrix fast-weight model whose entire recurrent state occupies 32,768 bytes answers queries at 0.9995 mean accuracy over three seeds, while a parameter-matched flat-vector recurrence and a parameter-matched transformer, trained on identical data for identical steps, both sit at chance. The paired-seed confidence intervals exclude the pre-registered 0.30 separation margin with lower bounds

above 0.95. The separation survives context extension: with the state fixed at 32,768 bytes, accuracy stays at or above 0.998 for every seed as the same episode is embedded at 454, 902, and 1,798 tokens, and the transformer stays at chance both uncapped and under every KV-cache byte budget from $1\times$ to $32\times$ the fast-weight state.

Because the baseline never demonstrates recall at any cache budget, the comparison licenses no “ $M\times$ less memory” headline. We follow the pre-registered analysis rule for this case and report the registered outcome verbatim: baseline non-competitive at matched params/tokens.

Our contributions:

- A matched three-arm separation on recall. Parameters matched within 2.8%, identical tokens, steps, optimizer, and paired episode seeds; the fast-weight arm clears the demonstration bar in every seed, the baselines never do (Section 3).
- A constant-memory demonstration. Recall accuracy is flat in context length out to eight times the binding span at a fixed 32,768-byte state, while no KV budget up to $32\times$ those bytes rescues the transformer (Section 4).
- A causal mechanism read. Zeroing the first block’s fast-weight state collapses recall to chance in every seed, the second block’s state proves inert, and no linear probe on either raw state decodes the bindings, so the storage is nonlinear and the model’s own forward pass is its decoder (Section 5).
- Scope, disclosed. A multi-hop variant of the task is a trainability finding, not a capability separation (3 of 9 contender seeds learn it; 0 of 9 baseline seeds), and at load $K/d = 0.75$ every arm reads chance (Section 6).

2 Experimental setup

2.1 Task: episode-restricted associative recall

Each training and evaluation sequence is an episode: $K = 32$ bind clauses followed by a query. A bind clause is a 7-token statement linking a key entity to a value entity, both drawn without replacement from a pool of real GPT-2 token identifiers (Radford et al., 2019), so the binding span is $T_{\text{bind}} = K \times 7 = 224$ tokens. The query names one key; the model must produce the bound value at the answer position. Keys and values are re-drawn every episode, so the map is in-context and cannot be memorized in slow weights. The primary task is single-hop; a multi-hop variant appears in Section 6.

Metric. We score episode-restricted recall accuracy: the argmax over the model’s own LM-head logits restricted to the episode’s K bound values, read at the answer position of the model’s own continuation. Chance is $1/K = 0.03125$, and an arm counts as demonstrating recall in a cell only above the pre-registered demonstration bar of $3/K = 0.09375$; readings at or below the bar are scored no-recall, and differences between two no-recall arms are treated as chance-level noise ordering, never as separation. One caveat travels with every accuracy in this paper: under argmax decoding, an associative state of rank one can already support on the order of d associations (Nichani et al., 2024), so recall accuracy here never establishes rank, continuous recovery, or storage capacity. Section 5 reports the continuous-readout diagnostics separately.

2.2 Three arms, matched budget

Matrix fast-weight model (the contender). A two-block decoder-style LM whose mixer holds a 64×64 matrix state S_t per block, updated by a delta rule (Schlag et al., 2021; Yang et al., 2024b): $S_t = S_{t-1}(I - \beta_t k_t k_t^\top) + \beta_t v_t k_t^\top$, with per-token read $o_t = S_t q_t$. Its total recurrent state is $2 \times 64 \times 64$ fp32 values, or 32,768 bytes, constant in context length.

Flat-vector recurrence (the additive control). Identical embedding table, output head, and feed-forward blocks; only the mixer changes, to a gated elementwise vector recurrence $s_t = s_{t-1} \odot g_t + v_t$ with read $o_t = s_t \odot q_t$, at matched total parameter count via mixer width

and $d_{\text{state}} = 64$ pinned equal to the contender’s. This arm removes the multiplicative outer-product update while preserving everything else, rather than reshaping the same operator, since any d^2 -dimensional vector can encode a $d \times d$ matrix and structure matters only if the operations use it. This control differs from the contender on three axes at once: mixer recurrence, raw state capacity (512 bytes versus 32,768, reported for completeness), and write conditioning. Its update has no dependence on the key k_t at all, so it cannot perform content-addressed binding under any training outcome. Vector-native schemes that do bind from a vector state use a key-dependent operation, circular convolution (Plate, 1995) or tensor-product binding (Smolensky, 1990); this control has none. Section 3 states which reading of the resulting comparison the capacity and write-conditioning gaps permit.

Transformer. A two-layer GPT-2-style causal decoder with the same tokenizer and width, evaluated in two roles: uncapped (its full KV cache), and KV-capped with attention-sink-plus-FIFO eviction (Xiao et al., 2023) at a total cache budget pinned to $M \times 32,768$ bytes across all layers and heads, swept over $M \in \{1, 2, 4, 8, 16, 32\}$.

Parameter counts are 14,049,408 (contender), 14,048,384 (flat-vector, -0.007%), and 14,440,448 (transformer, $+2.8\%$). All arms train from scratch for 20,000 steps at learning rate 3×10^{-4} with the same optimizer, batch schedule, auxiliary losses, and episode stream; sampling seeds are paired across arms, so the three arms of a given seed index see identical episodes.

2.3 Analysis protocol, frozen before the runs

Margins, tiers, and analysis rules were registered in the project’s experiment registry and frozen before the multi-seed runs. The primary comparison is contender versus flat-vector at the matched final checkpoint: a win requires the contender to clear the demonstration bar and the paired-seed 95% t -interval on the accuracy difference (three seeds, $\text{df} = 2$) to exclude a 0.30 margin. Two clauses matter for reading this paper. First, the registry pins a degenerate-baseline rule: if the uncapped transformer itself fails the demonstration bar at the primary cell, the capped-budget walk is definitionally degenerate, every cap is trivially cleared, and the registered outcome is “baseline non-competitive at matched params/tokens,” with no memory-multiplier tier certified. Second, a joint-no-recall rule scores any cell where both compared arms sit at the bar as a tie, never as a separation. Evaluation uses a pinned held-out query set of 4,096 queries over 128 episodes per cell, identical across arms and seeds. Three seeds is the registered count for the primary comparison; the protocol includes a seed-extension trigger for any cell whose paired gap approaches the 0.30 margin, and none did (Table 1), so the count was not extended for this comparison. The multi-hop variant’s later nine-seed table (Section 6) reflects a different, seed-variance-driven extension under the same protocol, not a change in standard. The protocol treats the primary contender-versus-additive-control comparison as the single pre-registered decision rule; the remaining reported intervals in this paper (the per- M capped-budget gaps, the pooled multi-hop interval, per-seed bar comparisons) are descriptive and diagnostic rather than additional hypothesis tests, so no multiple-comparison correction is applied to them.

3 Recall separation at matched budget

Table 1 reports the primary comparison. The contender reads 0.9995, 1.0000, and 0.9990 across the three seeds, at least $10\times$ the demonstration bar in every seed; the flat-vector arm reads 0.0322 to 0.0369 and the transformer 0.0271 to 0.0293, all within noise of chance and below the bar in every seed. The paired difference against the flat-vector arm is 0.9656 with a 95% interval of (0.9582, 0.9729); against the transformer, 0.9712 with (0.9685, 0.9738). Both intervals exclude the frozen 0.30 margin with more than 0.65 to spare at the interval floor, and no seed straddle fires the registered seed-extension trigger.

The readings this table supports are narrower than the raw gap suggests. Under the shared training budget, the matrix fast-weight update reaches ceiling while the additive control’s training loss barely moves past its first-500-step value for the remaining 97.5% of the run (Appendix C); learning rate, warmup, and auxiliary-loss weighting were shared across three

arm	seed 0	seed 1	seed 2	mean	clears bar?
matrix fast-weight	0.9995	1.0000	0.9990	0.9995	every seed
flat-vector recurrence	0.0322	0.0327	0.0369	0.0339	never
transformer (uncapped)	0.0271	0.0293	0.0286	0.0283	never
Δ vs flat-vector: mean 0.9656, 95% paired CI (0.9582, 0.9729); excludes 0.30					
Δ vs transformer: mean 0.9712, 95% paired CI (0.9685, 0.9738); excludes 0.30					

Table 1: Episode-restricted recall accuracy on the primary cell ($K = 32$, $K/d = 0.5$, binding span 224 tokens), per seed and mean, at the matched final checkpoint (20,000 steps, identical data and optimizer; chance $1/32 = 0.031$, demonstration bar $3/32 = 0.094$). The matrix fast-weight model clears the bar in every seed; neither the parameter-matched flat-vector recurrence nor the parameter-matched uncapped transformer ever does. Both paired-seed 95% t -intervals ($df = 2$) exclude the pre-registered 0.30 margin. Accuracy is argmax-decoded and episode-restricted, so it carries the capacity caveat of Section 2.

architecturally distinct mixers, not independently tuned per arm, so this paper does not claim the additive update is architecturally incapable of the task, only that it does not reach it under this budget. A second reading holds but is more modest than “matrix beats vector”: the additive control also holds $64\times$ less raw state and a write with no dependence on the query key (Section 2), so the comparison shows a matrix state with a multiplicative, key-conditioned write outperforms a smaller, non-key-conditioned vector state under a shared training budget; capacity and write-conditioning are not disentangled by this ablation, and no resized or key-conditioned vector control was run. The transformer column carries its registered phrasing, non-competitive at matched params/tokens: at this scale and budget it does not learn the task at all, which is itself the axis-2 datum, never extrapolated to differently-scaled or differently-trained attention models. What the table does not license is any claim that a transformer cannot perform this task, or that outer-product structure specifically, rather than capacity and key-conditioning together, is what the additive control’s failure isolates. Attention solves associative recall well in other regimes (Arora et al., 2023; Nichani et al., 2024); under this task distribution, tokenizer, parameter count, and step count, the attention arm did not get there while the fast-weight arm reached ceiling in every seed.

4 A fixed state versus a growing cache at long horizon

The constant-memory question is whether the recall of Section 3 survives when the bound episode sits inside a much longer context, with the fast-weight state’s byte size unchanged. We re-evaluate the same frozen checkpoints with each episode embedded in filler continuation at three horizons: 454, 902, and 1,798 tokens, that is, two, four, and eight times the binding span plus the query.

The fixed state does not decay. The contender reads 0.9995, 0.9983, and 0.9990 (per seed) at every horizon; the three curves in Figure 1 are flat to the third decimal from 454 to 1,798 tokens. The state occupies 32,768 bytes at every point on the curve.

No cache budget rescues the baseline. The uncapped transformer reads 0.029 to 0.036 across seeds and horizons; capped at every budget $M \times 32,768$ bytes for $M \in \{2, \dots, 32\}$ it reads 0.020 to 0.033, at or below its own uncapped reading everywhere. Per the frozen no-recall rule, these differences are chance-level orderings. Two registered hypotheses die here at once. A larger cache does not eventually hold the bindings: $M=32$ grants the cache more bytes than the episode needs and the reading does not move. Forced locality, which might have helped by concentrating attention on recent tokens, does not either; every capped reading sits at or below uncapped.

What is not claimed. The registered degenerate-baseline clause fires on this data exactly as written: the uncapped baseline itself fails the demonstration bar at the primary cell (readings 0.0271 to 0.0293), so every cap is trivially cleared and the per- M walk cannot certify a memory-multiplier tier. Although each of the five eligible per- M paired gaps is




Figure 1: Episode-restricted recall accuracy versus evaluation context length (log scale) on the primary cell. Each episode’s 224-token binding span is embedded in filler continuation out to 454, 902, and 1,798 tokens. The matrix fast-weight model (three seeds, solid) holds ≥ 0.998 at every horizon with its state fixed at 32,768 bytes. The transformer reads chance at every horizon, both uncapped (dashed) and under sink-plus-FIFO KV caps at $M \times 32,768$ bytes for $M \in \{2, \dots, 32\}$ (crosses; $M=1$ is excluded from the walk because its 8-token cap sits below the query length and is reported descriptively in Table 4). Reference lines mark chance ($1/32$) and the pre-registered demonstration bar ($3/32$).

individually clean (interval lower bounds all ≥ 0.9586 at the 902-token decision horizon), we quote no “the fast-weight model matches at $M \times$ less memory” figure, because there is no competent baseline reading to anchor the ratio. The registered outcome is the sentence used throughout this paper: baseline non-competitive at matched params/tokens. What the comparison yields instead is the contender’s flat horizon curve at constant bytes, together with the failure of every cap to change the baseline. Table 4 in the appendix lists the full grid.

5 Where the recall lives

A separation on accuracy alone leaves open where the winning arm keeps the bindings. We intervene directly: after the bind phase, the recurrent states are frozen, one block’s state is replaced with zeros, and the query continuation runs as a pure function of the remaining state and the query tokens. The decoder reads only the state, never the raw bind tokens, and this is checked rather than assumed: a fresh model instance given only the cached state (no raw bind-phase tokens) reproduces bit-identical continuation logits to the original run, closing the channel by which a hidden cache of the raw tokens could otherwise leak into the query pass.

The first block’s state is causally necessary; the second block’s is inert. Zeroing S_0 collapses per-seed accuracy from 0.9995, 1.0000, 0.9990 to 0.0339, 0.0012, 0.0002; zeroing S_1 leaves 0.9995, 0.9949, 0.9990. The same intervention on the flat-vector arm collapses its already-chance reading no further, and the pattern replicated in all twelve recurrent cells of the multi-seed run without a single localization failure. One instrument note: at seed 1 (intact accuracy exactly 1.0), S_1 -zeroing produces a small, reproducible change on 21 of 4,096 queries (0.51%, $\Delta = 0.0051$). The binomial standard-error band used elsewhere in this check assumes a sampling process and is ill-posed at $p = 1$ (it evaluates to zero width); both the intact and S_1 -zeroed readings are deterministic argmax decodes of the same frozen




Figure 2: Causal state-zeroing on the matrix fast-weight model, per seed on the primary cell (4,096 pinned queries per read; chance 1/32). Zeroing the first block’s fast-weight state S_0 before the query continuation collapses recall to chance; zeroing the second block’s state S_1 leaves it within noise of intact. The bindings are causally resident in S_0 .

checkpoint on the same fixed query set, so there is no sampling process for the band to characterize, the 21-query change is a real and repeatable effect of the intervention, and the band is not used to adjudicate this cell. The effect does not touch the S_0 -necessity result; whether those 21 queries reflect a small, consistent causal contribution from S_1 or near-tied logits is not established by this instrument.

The storage is nonlinear, and the model’s own forward pass is its decoder. Ridge probes fit offline on 24,576 fresh draws and evaluated on the pinned query set recover the bound value’s target vector from no linear read-out of either raw state: every (S, q) combine tested, including one placed directly on the causally necessary S_0 , reads within at most a +0.063 cosine gap of its shuffled control and passes a 0.9 cosine threshold on 0.0% of queries. Only the post-second-block, pre-LM-head hidden state is linearly decodable, and only in the fast-weight arm: cosine mean 0.894, threshold pass rate 0.674, gap +0.800 over its shuffled control. The flat-vector arm’s same site reads 0.119, pass rate 0.0. The bindings therefore live in S_0 in a format no tested linear map exposes; what transforms the S_0 -derived signal into a linearly legible code is the second block’s own forward computation, since its recurrent state is causally inert (Figure 2). Consistent with a nonlinear store read through a learned decoder, the pass rate of the pre-LM-head probe varies across seeds (0.686, 0.771, 0.951) while task accuracy is flat at ceiling; probe legibility is a property of the instrument. Appendix Table 5 lists all read-out sites.

6 Scope: what does not separate

Two further results bound the claim.

Multi-hop recall separates only as seed variance. On the multi-hop variant (compositional queries over a single Hamiltonian K -cycle, trained at hop depths 1–2, Table 2), the contender clears the demonstration bar in 3 of 9 seeds (0.334, 0.479, 0.391) and reads chance in the other six; the flat-vector arm reads chance in all nine. The registered adjudication rule scores this as trainability/seed-variance: the hypothesis that multi-hop recall is a hard capability boundary for this architecture at this scale is rejected (some seeds reach it), but the pooled paired interval $(-0.020, 0.268)$ straddles zero and the pre-registered batch-effect gate flags the pooled cohort (variance ratio 6.14 against a 4.0 bound), so no capability separation is claimed in either direction. The bar-clearing seeds also do not generalize: held-out hop depths 3–4 read chance, and the first bar-clearing seed’s partial recall collapses to 0.010

arm	s0	s1	s2	s3	s4	s5	s6	s7	s8
contender	0.040	0.029	0.334	0.036	0.031	0.479	0.028	0.391	0.041
flat-vector	0.034	0.028	0.037	0.031	0.031	0.032	0.034	0.031	0.035
pooled paired Δ CI (-0.020, 0.268), flagged non-decision-grade (var-ratio 6.14>4.0)									

Table 2: Multi-hop variant, nine seeds per compared arm, accuracy at the matched final checkpoint (chance 0.031, demonstration bar 0.094; readings above the bar in bold). Three of nine matrix fast-weight seeds learn the task; zero of nine flat-vector seeds do. The pooled paired interval straddles zero and is additionally flagged non-decision-grade by a pre-registered batch-effect gate, so this table is a trainability disclosure, not a capability claim.

	matrix fast-weight	flat-vector	transformer
$K=48$ ($K/d=0.75$) accuracy	0.0189	0.0195	0.0218
chance 0.0208; locate-only bar 0.0625; no arm clears			

Table 3: Stress cell at $K=48$ bindings on the same $d=64$ state ($K/d = 0.75$; chance 0.021). Every arm, including the matrix fast-weight model, reads chance: the recall separation of Table 1 is load-bounded and does not extend to this operating point at this scale and budget.

when its episode is embedded at any extended horizon. Multi-hop composition at this scale remains an open training problem.

The separation is load-bounded. At $K = 48$ bindings on the same 64-dimensional state ($K/d = 0.75$, Table 3), all three arms read chance: 0.0189 (contender), 0.0195 (flat-vector), 0.0218 (transformer). This cell was registered as locate-only, so it carries no tier arithmetic; what it establishes is that the ceiling-level recall at $K/d = 0.5$ does not come for free at higher load under the same training budget. Where between the two load points the capability degrades, and whether more training moves the boundary, are open questions this paper does not answer.

7 Related work

Fast-weight memory. Writing transient associations into multiplicative weights dates to Schmidhuber (1992), revived for attention-era models by Ba et al. (2016). Schlag et al. (2021) identified linear attention (Katharopoulos et al., 2020) as a fast-weight programmer and introduced the delta-rule update our contender uses; Yang et al. (2024b) made that rule trainable at scale. Selective state-space models (Gu & Dao, 2023) pursue the same constant-state property with a different recurrence, gated delta rules refine the contender’s own update (Yang et al., 2024a), and test-time memorization modules (Behrouz et al., 2024) share this paper’s motivating contrast between a fixed-size memory and a growing KV cache. None of these works reports a parameter- and token-matched comparison against both a transformer and a structure-removed vector control on episode-restricted recall, with a causal state intervention and a byte-pinned cache-cap protocol; that combination is this paper’s contribution.

Recall as the diagnostic. Arora et al. (2023) established multi-query associative recall as the axis separating attention from efficient alternatives, with the opposite headline: in their regime, attention wins and recurrent models lag. The two findings are consistent; ours sits at a different operating point: a small matched budget (14M-class models, 20,000 steps from scratch), higher load per state dimension ($K/d = 0.5$), an episode-restricted argmax metric, and a task distribution built for exact matched-budget control. We also report where our separation dies ($K/d = 0.75$, Section 6), which their framing predicts it must somewhere.

KV-cache reduction. Sink-plus-window eviction (Xiao et al., 2023) and heavy-hitter selection (Zhang et al., 2023) compress the cache of a model that already performs the task, asking how few bytes preserve its behavior. We use the sink-plus-FIFO mechanism of Xiao et al.

(2023) in the opposite role, as an evaluation instrument that pins the baseline’s cache to a byte budget matched against a fixed-size state, and our question is prior to theirs: whether the matched attention model learns the behavior at all.

Associative-memory capacity. Nichani et al. (2024) analyze factual recall in transformers via associative memories and prove that under argmax decoding a rank-one outer-product store supports on the order of d associations. We import their result as a standing caveat on what accuracy metrics can claim (Section 2); our contribution is empirical and comparative, not a capacity bound.

8 Limitations

The models are 14M-parameter class, trained 20,000 steps from scratch on a synthetic episode distribution with a GPT-2 tokenizer; nothing here speaks to pretrained models, natural text, or other scales. The baseline non-competitiveness holds at the matched budget only: a larger, longer-, or differently-trained transformer requires its own matched run before any extension of the claim. The metric is argmax-decoded and episode-restricted, so by the capacity results of Nichani et al. (2024) it supports no statement about rank, exact recovery, or how much the 32,768-byte state holds; the continuous diagnostics of Section 5 are reported only as bounds. The separation is demonstrated for one fast-weight family (a delta-rule matrix state) against one vector control and one transformer configuration; the multi-hop variant does not separate (Section 6); and the recall ceiling is load-bounded, with every arm at chance at $K/d = 0.75$. Horizon evaluation extends context with filler continuation, not with distractor bindings, which is the easier long-context condition and is disclosed as such. The longest tested context is 1,798 tokens, eight times the binding span: this paper demonstrates the constant-memory mechanism as a small-scale existence result, not a claim that it holds at the context lengths (tens of thousands to millions of tokens) where KV-cache cost is a practical deployment constraint.

9 Conclusion

Under a fully matched training budget, a two-block matrix fast-weight model stores and retrieves in-context bindings that neither matched baseline learns at all, from a state whose size does not grow with context. The bindings are causally localized to the first block’s fast-weight matrix and stored in a form only the model’s own forward pass decodes. A baseline that never becomes competent anchors no ratio, so no memory multiplier is on offer; the registered outcome of Section 4 is the paper’s only comparative claim. For the workshop’s question, whether long-horizon reasoning must pay a memory cost that grows with context, this is a small-scale existence result on the recall substrate such reasoning requires: at matched budget, the constant-memory architecture was the only one that learned to remember.

References

- Simran Arora, Sabri Eyuboglu, Aman Timalsina, Isys Johnson, Michael Poli, James Zou, Atri Rudra, and Christopher Ré. Zoology: Measuring and improving recall in efficient language models. arXiv preprint arXiv:2312.04927, 2023.
- Jimmy Ba, Geoffrey Hinton, Volodymyr Mnih, Joel Z. Leibo, and Catalin Ionescu. Using fast weights to attend to the recent past. arXiv preprint arXiv:1610.06258, 2016.
- Ali Behrouz, Peilin Zhong, and Vahab Mirrokni. Titans: Learning to memorize at test time. arXiv preprint arXiv:2501.00663, 2024.
- Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. arXiv preprint arXiv:2312.00752, 2023.
- Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. Transformers are RNNs: Fast autoregressive transformers with linear attention. arXiv preprint arXiv:2006.16236, 2020.

- Eshaan Nichani, Jason D. Lee, and Alberto Bietti. Understanding factual recall in transformers via associative memories. arXiv preprint arXiv:2412.06538, 2024.
- Tony A. Plate. Holographic reduced representations. IEEE Transactions on Neural Networks, 6(3):623–641, 1995. doi: 10.1109/72.377968.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI, 2019.
- Imanol Schlag, Kazuki Irie, and Jürgen Schmidhuber. Linear transformers are secretly fast weight programmers. arXiv preprint arXiv:2102.11174, 2021.
- Jürgen Schmidhuber. Learning to control fast-weight memories: An alternative to dynamic recurrent networks. Neural Computation, 4(1):131–139, 1992. doi: 10.1162/neco.1992.4.1.131.
- Paul Smolensky. Tensor product variable binding and the representation of symbolic structures in connectionist systems. Artificial Intelligence, 46(1-2):159–216, 1990. doi: 10.1016/0004-3702(90)90007-M.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. arXiv preprint arXiv:1706.03762, 2017.
- Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. Efficient streaming language models with attention sinks. arXiv preprint arXiv:2309.17453, 2023.
- Songlin Yang, Jan Kautz, and Ali Hatamizadeh. Gated delta networks: Improving Mamba2 with delta rule. arXiv preprint arXiv:2412.06464, 2024a.
- Songlin Yang, Bailin Wang, Yu Zhang, Yikang Shen, and Yoon Kim. Parallelizing linear transformers with the delta rule over sequence length. arXiv preprint arXiv:2406.06484, 2024b.
- Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuandong Tian, Christopher Ré, Clark Barrett, Zhangyang Wang, and Beidi Chen. H₂o: Heavy-hitter oracle for efficient generative inference of large language models. arXiv preprint arXiv:2306.14048, 2023.

A Full capped-budget grid

arm / KV budget	$T=454$	$T=902$	$T=1798$
contender, fixed 32,768 B	0.9989	0.9989	0.9989
transformer, uncapped	0.0293	0.0314	0.0339
transformer, $M=32$ cap	0.0300	0.0251	0.0256
transformer, $M=16$ cap	0.0280	0.0240	0.0251
transformer, $M=8$ cap	0.0257	0.0264	0.0250
transformer, $M=4$ cap	0.0251	0.0277	0.0251
transformer, $M=2$ cap	0.0251	0.0286	0.0257
per- M paired-gap CI lower bounds at $T=902$: all ≥ 0.9586			

Table 4: Mean episode-restricted recall accuracy (three seeds) by arm, KV budget, and evaluation context length on the primary cell. The matrix fast-weight model’s state is 32,768 bytes at every entry of its row; the transformer’s cap at budget M pins its total KV bytes across layers and heads to $M \times 32,768$. $M=1$ is excluded from the registered walk because its cap length (8 tokens) sits below the query length; its descriptive readings, 0.021 to 0.030 across seeds and horizons, sit at chance like every other transformer entry. Chance is 0.031; the demonstration bar is 0.094; no transformer entry approaches either.

B Linear read-out sites

arm	linear read-out site	cos. mean	rf@0.9	shuffled cos.	gap
contender	S_1q (shallow query)	0.165	0.000	0.105	+0.060
contender	S_0q (bind-path query)	0.118	0.000	0.112	+0.006
contender	S_1q (true query)	0.167	0.000	0.104	+0.063
contender	pre-LM-head hidden	0.894	0.674	0.094	+0.800
flat-vector	S_1q (shallow query)	0.117	0.000	0.112	+0.005
flat-vector	S_0q (bind-path query)	0.113	0.000	0.111	+0.002
flat-vector	S_1q (true query)	0.118	0.000	0.113	+0.005
flat-vector	pre-LM-head hidden	0.119	0.000	0.112	+0.006

Table 5: Offline ridge-probe recovery of the bound value’s target vector by read-out site, on the single-seed localization checkpoints (fit on 24,576 fresh draws, evaluated on the pinned 4,096-query set; “rf@0.9” is the fraction of queries whose predicted vector clears cosine 0.9 against the target; “gap” is cosine mean minus a shuffled-fit control). No linear combine of a raw recurrent state and the query clears the threshold in either arm, including the read placed directly on the causally necessary S_0 ; only the pre-LM-head hidden state decodes, and only in the matrix fast-weight arm.

C Training-loss trajectories, all three arms




Figure 3: Training loss (primary task, seed 0) over the full 20,000-step run, all three arms, from the same archived configs used for the parameter counts in Section 2. The flat-vector and transformer arms drop to 7.83 and 7.84 by step 500 and move only to 7.69 and 7.51 over the remaining 19,500 steps; the matrix fast-weight arm keeps falling for the entire run, reaching 1.38. Hyperparameters are the shared settings of Section 2; no per-architecture learning-rate sensitivity sweep was run, and Section 3 scopes its claim to this shared budget accordingly.

D Reproducibility of every number

Every figure and every table cell in this paper regenerates from archived evaluation artifacts by a single versioned script that asserts the MD5 checksum of each source file before loading it; a changed artifact fails the build rather than silently changing a number. Every inline numerical claim in the source carries a machine-checkable comment naming its row in a claims-to-evidence map maintained alongside the draft. The archived artifacts, the script, and the map will be released with the de-anonymized version of this paper.