

When the Gradient Sees Rank: Provable Necessity, Causal Recruitment, and Exact Composition in Trained Matrix Memories

Editors: Under Review for the NeurReps Workshop

Abstract

Matrix-valued fast-weight memories make rank a structural budget: the directions a state spans bound what it can bind and compose. A recent negative result reported that a matrix latent bolted onto a vector-pretrained reasoner never uses its rank, but the probed task admits a rank-1 solution, leaving the general question open. We answer it on testbeds where the exact solution provably requires $\text{rank}(Z) \geq K$ and where every rank-evading shortcut is closed by construction: the readout is pinned to exact continuous recovery, since argmax decoding lets a rank- m memory store roughly md associations; a single-matrix-state bottleneck forecloses position decomposition; and no cell is called dead without re-testing at $2\text{--}2.5\times$ budget. Under these conditions gradient descent recruits the rank the task demands in this architecture family: learned effective rank tracks K across the tested grid (Spearman $\rho = 1.0$ at $d = 16$), train-time rank caps show a causal step at $k \approx K$ (at $d = 8$, $K = 4$: rank 3 gives at most 0.0004 recovery, rank 4 gives 0.97), and the trained operator composes exactly through 21-fold self-application in four of five seeds. An entity-subspace decomposition supplies the mechanism, restricted rank exactly K on the key subspace, and in the single converged rank-starved seed, the depth-decay curve is predicted from its eigenspectrum alone.

Keywords: effective rank, fast weights, associative memory, composition

1. Introduction

Matrix-valued latents, in continuous chain-of-thought models (Hao et al., 2024) and fast-weight architectures broadly (Schlag et al., 2021), carry a structural observable a vector does not have: the *rank* of the state, which should track how many items the state holds and, when constrained below what the task needs, destroy performance. A recent negative result tested this on a matrix-CODI model (a $d \times d$ latent bolted onto a vector-pretrained GPT-2) and found rank- k truncation curves flat across four training regimes and four readouts (Larson, 2026); its own stated confound is that ProsQA admits a rank-1 solution, so flatness says nothing about tasks that provably require rank. We close the confound: matrix-native models trained from scratch on tasks whose exact solutions carry a provable $\text{rank}(Z) \geq K$ lower bound, under a readout and architecture built so nothing short of genuine rank recruitment can pass. The contributions: the *teeth* that keep the bound a necessity result (§2); recruitment of rank tracking K and its causal necessity under train-time caps (§3); exact composition to depth 21, its invariant-subspace mechanism, and the depth-amplification signature (§4); and an exactness frontier in d with a budget-extension rule (§5).

2. The Bound and the Teeth That Keep It a Bound

Task and bound. Each episode presents K key→value bindings, then queries one; keys and values are fresh per episode (Gaussian, normalized, or exactly orthonormal following Mezzadri, 2007), so evaluation uses never-seen bindings. Exact recovery $Zk_j = v_j$ for K bindings with linearly independent keys and values forces $ZK_{\text{mat}} = V_{\text{mat}}$, hence the classical bound $K = \text{rank}(V_{\text{mat}}) = \text{rank}(ZK_{\text{mat}}) \leq \text{rank}(Z)$ (Kohonen, 1972; Anderson, 1972). We measure whether a trained, bottlenecked model *discovers* this rank under plain SGD, and whether it is causally load-bearing.

Tooth 1: exact continuous recovery, never argmax. Nichani et al. (2025) construct a rank- m linear memory storing approximately md associations under interference-tolerant argmax decoding; under that rule a rank-1 state passes at every K we test and the bound collapses. We therefore pin the decoder to the linear unbind itself, $\text{pred} = Z \cdot k_j$, scored by an absolute bar (rec@0.9: the fraction of queries with $\cos > 0.9$), never an argmax and never a learned readout that could launder rank.

Tooth 2: a single-state bottleneck, verified behaviorally. “Hold K items” is satisfiable in full attention by attending to K raw-input positions at rank 1 each, so all information funnels through one $d \times d$ state and the decoder is a pure function of (Z, query) : after the encoder writes Z , the raw-input cache is corrupted and the decode is confirmed bit-for-bit unchanged (a gradient blank-out test, not a mask-shape check).

Model and controls. A Transformer encoder over the binding tokens, with d learned row-reader latents producing the rows of $Z \in \mathbb{R}^{d \times d}$: permutation-invariant, able to express any rank up to d , and not hard-coding $\sum_j v_j k_j^\top$. Every model is under 1M parameters (the $d = 16$ composition model: 170,896). Effective rank (entropy of the normalized singular-value spectrum) is the pre-registered primary metric; the primary causal control constrains Z to rank k at *every training step* (spectral projection), not post hoc.

3. Rank Tracks K and Is Causally Necessary

Recruitment. Without any rank term in the loss, the trained $d=16$ state’s effective rank rises one-for-one with binding count across a ten-point grid from $K=1$ to $K=16$ (Spearman $\rho = 1.0$; 2.42 at $K=1$, 8.20 at $K=8$, 15.09 at $K=16$), leaving the pre-registered $[0.7K, 1.3K]$ band only at $K=1, 2$, by overshoot; beyond $K = d$ rank saturates and declines. A subsequent, larger replication sweep is directionally identical.

Causal necessity. Training with Z constrained to rank k throughout (Figure 1): at $d=8, K=4$ the step is exact, $k \leq 3$ yields $\text{rec@0.9} \leq 0.0004$ and $k = 4$ yields 0.97, a discontinuity at the provable bound. The $d=16$ cells show the same knee near $k \approx K$ with mild post-knee non-monotonicity from seed noise; the replication resolves the $K=12$ transition into a ragged ramp, not a clean step. The two halves answer Larson (2026): the earlier rank-blindness was a property of a rank-1-solvable task, not of the gradient.

4. Exact Composition and the Invariant-Subspace Mechanism

Binding is a single lookup; the composition variant relates the K entities by one Hamiltonian K -cycle π , so a hop-depth- h query means recovering $v_{\pi^h(i)}$ by h -fold self-application of the same trained operator, $Z^h k_i$, with no learned per-hop parameters. Training sees hops




Figure 1: The primary, pre-registered causal test: rec@0.9 under train-time force-rank k , regenerated from the archived sweep (dashed line at $k = K$). The step lands at $k \approx K$ in every cell whose unconstrained model reaches exact recovery.




Figure 2: Depth as a spectral-exactness amplifier ($d=16$, $K=8$). *Left*: mean cosine for the converged force-rank $k=7=K-1$ seed against the eigenspectrum-only prediction (no raw keys). *Right*: rec@0.9 by hop for the four converged rank-sufficient seeds (flat at 1.0), the rank-starved operator, and the still-transitioning seed.

$\{1, 2, 3\}$; evaluation adds held-out $\{4, 5, 6\}$ and probes $\{7, 21\}$ ($d = 16$, $K = 8$, orthonormal keys). A single full cycle is load-bearing: under a general random permutation, nominally held-out depths collapse into trivial or in-distribution queries modulo its short cycle lengths (Appendix A).

Exact composition. 4/5 unconstrained seeds reach $\text{rec@0.9} \geq 0.9996$ at every tested hop (3/5 at 1.000 throughout); the fifth is still transitioning at budget, decaying from 0.93 ($h=1$) to 0.16 ($h=7$) and 0.0001 ($h=21$).

Depth amplification. At force-rank $k = K - 1$ (provably insufficient), the one converged seed passes shallow hops at mean cosine 0.92–0.93 ($h \leq 5$), above the 0.9 bar, and collapses only under depth: rec@0.9 falls $0.996 \rightarrow 0.881 \rightarrow 0.060$ at $h = 1/7/21$. The mechanism is spectral and predictable: from the operator’s eigenspectrum alone, predicted cosine matches the measured curve within 0.008 through $h = 7$ (Table 1).

The invariant-subspace mechanism. Whole-matrix effective rank of the converged solutions reads the full ambient dimension (16.0 of $d=16$; stable rank 14.6–15.6); the whole-matrix number is the wrong instrument. Block-decomposing the trained Z against the ideal cycle’s K -dimensional entity subspace $\mathcal{E}$: the restricted operator has effective rank 7.999–8.000 (maximum 8), matching the ideal scaled cycle to a scale-corrected residual of 0.7–2.4%; leakage out of $\mathcal{E}$ is below 1.5% of the state norm; the complement block is full-rank and non-contractive yet dynamically invisible: a query starting in $\mathcal{E}$ never leaves it

(Appendix B). Restricted to the subspace the task uses, “rank tracks K ” holds exactly; dead force-rank seeds fail the same restricted test, so their collapse is real.

5. The Exactness Frontier, and When to Trust a Dead Cell

Budget reversals. Three independent axes each declared a cell dead at its first step budget before an extension reversed the verdict: composition at $K=8$ (1/5 seeds converged at 20K steps; 4/5 at 40K), composition at $K \in \{12, 16\}$ (0/5 at 40K; 3/3 and 2/3 fresh seeds at 80K), and d -scaling (an 8K-budget wall at $d \geq 32$ dissolving at 20K, onset 6–16K). The rule, *a cell is not dead until re-tested at 2–2.5 $\times$ budget*, is necessary but not sufficient: $K=16$ ’s remaining stuck seed stayed dead through 120K steps, and the $d=32$ cells below plateau flat yet still fail the pre-registered bar.

The frontier that survives is not a trainability wall but an *exactness* frontier. Fixing $K = d/4$ at encoder width $h = 64$ (every cell confirmed flat before being read as converged), mean recovery cosine falls monotonically across d : 0.877/0.909/0.915 at $d=32$ (three seeds, 100K steps), 0.7196 at $d=48$ (100K), and 0.3882 at $d=64$ (150K); the exact-match metric is harsher, with $\text{rec}@0.9}$ 0.413/0.632/0.653 at $d=32$ (all below the pre-registered 0.7 pass bar), 0.002 at $d=48$, and 0.0 at $d=64$. Rank recruitment and exactness separate: at $d=32$ recruited rank reaches 91–97% of target while exact recovery stalls below the bar.

6. Related Work, Limitations, and Outlook

Nichani et al. (2025) is the nearest neighbor: a hand-built existence construction under interference-tolerant argmax decoding; this work pins the readout so that construction cannot pass, then measures and intervenes on the rank SGD recruits (Barnfield et al., 2026 give static-memory thresholds). Nazari and Rusch (2026) and Sun et al. (2026) measure effective-rank dynamics in pretrained linear-attention models observationally (an *upper* bound $\text{rank}(S_t) \leq t$); here the write count is controlled and necessity is causal. Mishra et al. (2026) train matrix-state RNNs on S_3 with no rank measurement or intervention (their own vector-state control matches); Grazzi et al. (2025) and Siems et al. (2025) (the Householder extension of Yang et al., 2024) characterize the per-token *transition* operator, not the trained *state*’s own rank, the quantity here. Continuous-CoT capacity results are existence constructions, not measurements (Zhu et al., 2025; Gozeten et al., 2025); the compositional-generalization caution (Dziri et al., 2023; Wang et al., 2024) is met with an exact eigenstructure match.

Limitations. Every experiment is synthetic and under 1M parameters; generality at scale, and across state-writing architectures (recurrent fast-weight, attention-readout, associative-memory), is untested. The depth-amplification contrast rests on the single converged force-rank- $(K-1)$ seed; the other two at that cell collapsed under a documented numerical instability in the spectral-projection backward pass, so it is a case study consistent with the spectral prediction, not an $n>1$ mechanism; the $d \geq 48$ frontier cells rest on single seeds. The registered falsifiers stand (Appendix C). A concurrent companion submission (group-composition state tracking) shares no figures or tables but restates this paper’s recruitment, causal-necessity, and composition numbers as background for its distinct contribution, a group-dimension rank law this paper does not make.

Appendix A. Depth 21 Under the Single K -Cycle

Under a single 8-cycle, $\pi^{21} = \pi^{21 \bmod 8} = \pi^5$, so the nominal depth-21 probe shares its *target* with the already-tested depth 5; the two are not distinct group-theoretic queries. What depth 21 measures is numerical stability under 21 sequential self-applications of the trained operator. In the four converged seeds rec@0.9 reads ≥ 0.9996 at both $h=5$ and $h=21$, so repeated self-application does not amplify their residual error; in the rank-starved operator the same 16 additional applications compound a passing shallow-hop cosine into collapse (0.947 at $h=5$ to 0.060 at $h=21$ in recovered fraction), which is exactly the amplification claim of §4, not a generalization claim to unseen permutation targets.

h	predicted cos	measured $\overline{\text{cos}}$	measured rec@0.9
1	0.9317	0.9303	0.996
3	0.9258	0.9259	0.986
5	0.9181	0.9212	0.947
7	0.9089	0.9163	0.881
21	0.7536	0.8206	0.060

Table 1: The rank-starved ($k = 7 = K - 1$) operator’s depth-decay curve: cosine predicted from the trained operator’s own eigenspectrum (no raw keys) against the measured values, with the recovered fraction alongside. Every hop through $h = 7$ matches within 0.008; at $h = 21$ the per-item distribution widens and the exact-match metric collapses while mean cosine moves little.

Appendix B. Per-Seed Subspace Decomposition

seed	effrank(A)	residual	leakage (% of $\ Z\ $)	effrank(D)	$\rho(D)$
s1	8.000	0.7%	0.44%	8.000	1.38
s2	7.999	1.9%	1.06%	8.000	1.27
s3	7.999	1.5%	0.87%	8.000	2.86
s4	7.999	2.4%	1.45%	7.999	1.02

Table 2: Entity-subspace block decomposition of the four converged $K=8$ composition seeds (means over four evaluation episodes), recomputed from the archived Z -dumps. A is the operator restricted to the K -dimensional entity subspace; residual is the scale-corrected Frobenius distance to the ideal cycle after fitting the single isotropic scale that cosine scoring cannot see; leakage is the cross-block norm $\|C\|$ as a fraction of $\|Z\|$; D is the complement block, full-rank and non-contractive ($\rho(D) \geq 1$) yet never visited by a query trajectory that starts in the entity subspace.

Appendix C. Reproducibility

Training, evaluation, and analysis code is available at <https://anonymous.4open.science/> (anonymized for review). Raw per-run JSON archives (dense training-step checkpoints, not summary rollups) back every number in this abstract; each figure regenerates from those archives through a single versioned script that asserts an md5 checksum for every source file it loads, so a changed artifact fails the build rather than silently plotting stale data. Hypotheses, metrics, force-rank controls, and decision bands were registered in design documents before training, and every deviation (one metric re-registration, logged before the affected sweep) is recorded inline in those documents. The registered falsifiers stand open: force-rank- $(K-1)$ at $\geq 0.9\times$ ceiling at multiple (d, K) points, or a $d \geq 32$ cell clearing the 0.7 bar, would overturn the causal claim. The force-rank staircase aggregate retains one recovery figure per (d, K, k) cell, not per-seed instability diagnostics; the necessity direction (recovery ≈ 0 below K) is robust to the cause of any individual sub- K collapse, while the smoothness of the post-knee approach should be read at the aggregate level only. The analysis pipeline emits benign BLAS RuntimeWarnings from near-singular intermediate products on some platforms; every decomposition output is asserted finite before use, so these do not affect any reported number.

References

- James A. Anderson. A simple neural network generating an interactive memory. *Mathematical Biosciences*, 14(3–4):197–220, 1972.
- Nicholas Barnfield, Juno Kim, Eshaan Nichani, Jason D. Lee, and Yue M. Lu. Sharp capacity thresholds in linear associative memory: From winner-take-all to listwise retrieval. *arXiv preprint arXiv:2605.05189*, 2026.
- Nouha Dziri, Ximing Lu, Melanie Sclar, Xiang Lorraine Li, Liwei Jiang, Bill Yuchen Lin, Peter West, Chandra Bhagavatula, Ronan Le Bras, Jena D. Hwang, Soumya Sanyal, Sean Welleck, Xiang Ren, Allyson Ettinger, Zaid Harchaoui, and Yejin Choi. Faith and fate: Limits of transformers on compositionality. *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. arXiv:2305.18654.
- Halil Alperen Gozeten, M. Emrullah Ildiz, Xuechen Zhang, Hrayr Harutyunyan, Ankit Singh Rawat, and Samet Oymak. Continuous chain of thought enables parallel exploration and reasoning. *arXiv preprint arXiv:2505.23648*, 2025.
- Riccardo Grazi, Julien Siems, Arber Zela, Jörg K. H. Franke, Frank Hutter, and Massimiliano Pontil. Unlocking state-tracking in linear RNNs through negative eigenvalues. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2411.12537.
- Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian. Training large language models to reason in a continuous latent space. *arXiv preprint arXiv:2412.06769*, 2024.
- Teuvo Kohonen. Correlation matrix memories. *IEEE Transactions on Computers*, C-21(4): 353–359, 1972.

Sam Larson. The gradient does not see rank: Rank-indifference in matrix-CODI on ProsQA. In *ICML 2026 Workshop on Mechanistic Interpretability*, 2026.

Francesco Mezzadri. How to generate random matrices from the classical compact groups. *Notices of the American Mathematical Society*, 54(5):592–604, 2007. arXiv:math-ph/0609050.

Mayank Mishra, Shawn Tan, Ion Stoica, Joseph Gonzalez, and Tri Dao. M²RNN: Non-linear RNNs with matrix-valued states for scalable language modeling. *arXiv preprint arXiv:2603.14360*, 2026.

Philipp Nazari and T. Konstantin Rusch. The key to state reduction in linear attention: A rank-based perspective. *arXiv preprint arXiv:2602.04852*, 2026.

Eshaan Nichani, Jason D. Lee, and Alberto Bietti. Understanding factual recall in transformers via associative memories. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2412.06538.

Imanol Schlag, Kazuki Irie, and Jürgen Schmidhuber. Linear transformers are secretly fast weight programmers. In *International Conference on Machine Learning (ICML)*, 2021. arXiv:2102.11174.

Julien Siems, Timur Carstensen, Arber Zela, Frank Hutter, Massimiliano Pontil, and Riccardo Grazi. DeltaProduct: Improving state-tracking in linear RNNs via householder products. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2502.10297.

Ao Sun, Hongtao Zhang, Heng Zhou, Yixuan Ma, Yiran Qin, Tongrui Su, Yan Liu, Zhanyu Ma, Jun Xu, Jiuchong Gao, Jinghua Hao, and Renqing He. State rank dynamics in linear attention LLMs. *arXiv preprint arXiv:2602.02195*, 2026.

Boshi Wang, Xiang Yue, Yu Su, and Huan Sun. Grokked transformers are implicit reasoners: A mechanistic journey to the edge of generalization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2405.15071.

Songlin Yang, Bailin Wang, Yu Zhang, Yikang Shen, and Yoon Kim. Parallelizing linear transformers with the delta rule over sequence length. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2406.06484.

Hanlin Zhu, Shibo Hao, Zhiting Hu, Jiantao Jiao, Stuart Russell, and Yuandong Tian. Reasoning by superposition: A theoretical perspective on chain of continuous thought. *arXiv preprint arXiv:2505.12514*, 2025.