Three Bounds on a Null: Testing the Link Between Fast-Weight Write Geometry and In-Context Composition

Anonymous authors
Paper under double-blind review

Abstract

Fixed-state sequence models bind in-context associations in a matrix-valued fast-weight state, and geometric observables of that state, such as the span occupied by key writes, have been proposed as lenses on what such models can compose. We test whether one such observable predicts, or causally improves, in-context multi-hop composition in DeltaNet-family language models spanning 14M to 1.31B parameters, and we find a null bounded three ways. First, a state-space composition readout returns a recovered fraction of exactly zero at 366 of 366 readings across three structurally different instruments (zero-shot marker prompts, natural language, and task-familiarized checkpoints); the readout's pre-registered validity gates fail at every cell, so the zeros are an instrument verdict, not a refutation. Second, a vocabulary-space behavioral contrast bounds any causal effect of a write-geometry intervention below its pre-registered three-seed detection floor of 1.5 to 1.7 loss units; the single determinate signal is a transient deviation in the harm direction. Third, a targeted twelve-seed replication of that transient was stopped by its own pre-registered batch-effect gate (cohort variance ratio 4.47 against a pinned cutoff of 4.0), and the new cohort's interval spans zero. Each bound names the observation that would overturn it. The complete program consumed approximately 9.5 GPU-hours.

1 Introduction

Linear-attention and delta-rule language models compress an entire context into a fixed-size matrix state (Schlag et al., 2021; Yang et al., 2024). A companion research program measured geometric observables of the written state (the span fraction occupied by key writes, the effective rank of the key population) across scales and under targeted interventions. Those measurements carried an unexamined promise: that write geometry matters for something a user would recognize, such as composing multi-hop relations presented in context. This paper reports the program that tested the promise; the test produced a null. A reader of a null has two standing worries, that the instrument never worked and that the analysis could declare a null after seeing the data. Three pre-registered bounds address both:

1. An instrument bound. A state-space composition readout, deployed three structurally different ways, returns a recovered fraction of exactly zero at 366 of 366 readings on checkpoints from 14M to 1.31B parameters, and its pre-registered validity gates fail at every cell, routing the outcome to probe-invalid, not refutation. One nonzero reading or one gate pass would have voided this bound.
2. A power bound. A vocabulary-space behavioral contrast, with a three-seed detection floor of 1.5 to 1.7 loss units registered before unblinding, leaves three of four contrasts unresolved; its single determinate signal is a transient deviation in the harm direction. A determinate effect above the floor would have voided this bound.
3. A replication bound. A targeted twelve-seed extension of that transient was stopped by its own pre-registered batch-effect gate (a 4.47-fold cohort variance mismatch against a pinned 4.0), and the new cohort's interval spans zero. A clean pooled interval excluding zero would have voided this bound.

Wave (instrument)	Grid	Readings	Registered outcome
1 zero-shot, marker	78 cells, 14M–392M	$312/312 = 0.0$	probe-invalid
2 zero-shot, 1.31B rung	2 cells	$8/8 = 0.0$	probe-invalid
3 zero-shot, natural lang.	4 gate cells	$16/16 = 0.0$	probe-invalid, refused
4 familiarized gate	6 cells × 5 ckpts	$30/30 = 0.0$; L_q falls	dissociation
5 behavioral contrast	18 cells, 4 contrasts	3 unresolved, 1 transient	power-bounded
6 seed extension, n=12	18 pretrain cells	variance ratio 4.47	batch-effect-flagged

Table 1: The program at a glance. Waves 1–4 deploy the geometric readout (recovered fraction at cosine ≥ 0.9); its validity gates fail at every cell of every wave, so all zero readings route to probe-invalid rather than refutation. Waves 5–6 deploy the vocabulary-space contrast; the one determinate transient is in the harm direction and its n=12 replication was stopped by the pre-registered batch-effect gate.

The contribution is the bounded null and the measurement discipline that kept it interpretable.

2 Models, Task, and Instruments

Checkpoints and task. All experiments probe DeltaNet-family language models (Yang et al., 2024) pretrained on a reasoning-dense corpus mix (OpenR1-Math) and an encyclopedic mix (WikiText-103, Merity et al., 2016). An intervention grid (14M parameters) varies a frozen-key-bias intervention on the fast-weight write path in three validation-loss-matched arms: control, per-token, and global. A scale ladder spans 14M, 98M, 392M, and 1.31B parameters (state dimension 64, 64, 128, 128; depth 2, 12, 16, 22). Each probe episode binds K entities into one directed cycle with bind statements, then queries the entity $h \in \{1, \ldots, 4\}$ hops around the cycle; a full K-cycle prevents held-out hop depths from collapsing modulo short cycle lengths. Load is K relative to state dimension ($K \in \{20, 32\}$ on the grid; near-capacity per rung).

Instrument 1: a geometric readout. For a layer's terminal state S_T and effective query q_{eff}, the readout predicts the h-hop target as $S_T^h q_{\text{eff}}$ and scores exact continuous recovery: the fraction of queries reaching cosine at least 0.9 with the target value vector. Continuous recovery, rather than argmax over candidates, follows Nichani et al. (2024), under which argmax decoding can succeed without the state carrying the claimed structure. Registered validity gates accompany the readout: an h=1 sanity floor (one-hop recovery must beat a shuffled-pairing null and an absolute 0.10 floor) and two alignment premises, (iii) query–key and (iv) key–value, each requiring a same-entity median cosine above a shuffled-entity null's 95th percentile. Gate failure routes to probe-invalid, never to refutation.

Instrument 2: a behavioral contrast. A familiarization phase trains checkpoints for 5,000 steps on the bind/query task mixed into the pretraining corpus; the contrast is the held-out-episode query loss L_q (a vocabulary-space cross-entropy through the language-model head, machinery disjoint from Instrument 1), differenced control minus arm at five checkpoints. A six-outcome trajectory classification and a differential condition (the K=32 effect determinate while K=20 is not) were fixed before unblinding. Every wave ran gate-first, with verdicts routed mechanically (Appendix A reproduces all gates and routing rules).

3 The Geometric Readout Never Fires

Zero-shot, marker template (wave 1). The full pre-registered grid, 60 intervention cells and 18 ladder cells (Table 1), returns a recovered fraction of exactly 0.0 at all 312 (cell, hop) readings. The validity gates fail everywhere: the h=1 floor fails both conditions, and premises (iii) and (iv) pass at 0 of 78 cells each. Figure 1 (Appendix B) shows the structure: premise medians track their shuffled nulls but never cross them (premise (iii) medians span $[-0.33, 0.76]$), and prediction–target cosines center near zero, ten-plus standard deviations from the 0.9 threshold. The registered routing reports probe-invalid; no refutation is licensed.

Scale (wave 2). The 1.31B rung reproduces the failure in kind: 8 of 8 readings at 0.0, both premise gates failing at both corpora, for a ladder series of 80 of 80 zeros across 14M to 1.31B: not a small-model artifact.

Zero-shot, natural language (wave 3). Two template families, a gift-verb family with the out-of-distribution query marker dropped and a corpus-native succession-verb family, fail identically: 16 of 16 readings at 0.0, the h=1 gate false at all 4 cells. The failure is not attributable to prompt surface form.

Task-familiarized (wave 4). The strongest-case test trains control checkpoints on the task itself for 5,000 steps and applies the gate at five checkpoints per cell. It fails at 30 of 30 readings, 0 of 512 queries recovered each time, while the vocabulary-space query loss falls 21.8 to 46.4 percent (mean 35.9 percent) in 6 of 6 cells (Figure 2, Appendix B): the models measurably learn the task through one readout while the other stays at floor. Since the two share no machinery by construction, the dissociation indicts the readout construct, not the models (the fall stops short of a registered 50-percent pin, so the adjudication claims partial task learning only).

Gates with enforcing code paths. Wave 1's launch chain computed its abort gate but had no code path that acted on it; the grid ran to completion, spending roughly 0.29 GPU-hours past a gate that had already declared the probe invalid. Every later chain enforced its gates at the process boundary, and waves 3 and 4 each refused a full-grid launch mechanically on a real failing gate. A registered gate is a claim about behavior; only an enforcing code path makes it one about the system.

4 The Behavioral Contrast Is Power-Bounded

Wave 5 promotes the vocabulary-space contrast to the primary instrument: all three arms are familiarized under identical recipes (18 cells), and the arm effect is the held-out query-loss difference, control minus arm, so positive means the intervention helps.

The registered power floor. Before unblinding, the three-seed detection floor was derived from the control cells' between-seed spread: a σ proxy of 0.43 to 0.48 loss units gives a 95-percent half-width of roughly 1.5 to 1.7. For calibration, the entire familiarization effect is about 1.69 loss units: the instrument can detect only an arm effect about as large as everything the model learns in the whole phase. The registration says in advance that unresolved is the most likely outcome for any modest effect, to be read as a measurement bound, never as evidence of no effect.

Results. Three of four (corpus, arm) contrasts classify as unresolved. The fourth (encyclopedic corpus, per-token arm) classifies as transient: at checkpoint 2,500 the K=32 effect is determinate at -0.500, interval $[-0.624, -0.376]$, while K=20 is not (mean -0.252, interval $[-0.920, +0.416]$), and by checkpoint 5,000 it dissolves (mean -0.795, interval $[-2.513, +0.923]$) (Figure 3, left, Appendix B). A held-out-hop secondary readout fires at the same cell and checkpoint, same direction, consistent with one coherent mid-training deviation.

Sign discipline. The intervention arm's loss sits above control: the one determinate signal points in the harm direction, and every determinate high-load reading in the primary table is negative. Two registered-forbidden misreadings loom: "the intervention matters" ignores that transients were registered as training-dynamics artifacts; "no effect" ignores the power floor. The wave delivers a bound plus one real, bounded, harm-direction transient.

5 The Replication Gate Refuses the Pool

Wave 6 extends the transient's anchor cell (encyclopedic corpus, per-token arm, K=32, checkpoint 2,500) to twelve seeds, with three integrity mechanisms registered before launch: nine fresh pretraining seeds for both arms (18 new cells); an archived-values loader with a runtime guard so the original three seeds are never re-scored live (their deltas appear byte-identical in the new harvest); and a batch-effect gate requiring, before any pooling, cohort means within twice the pooled standard error and a variance ratio below a pinned 4.0.

The gate fired. The archived cohort's between-seed standard deviation (1.154) is 2.1 times the new cohort's (0.546): variance ratio 4.47, twelve percent past the cutoff, while the means agree comfortably (shift 0.364 against a 1.382 threshold). The registered routing reports the cohorts separately and classifies the wave batch-effect-flagged, with no decision-grade pooled interval. The archived cohort reproduces its original interval $[-0.624, -0.376]$ by construction; the new nine-seed cohort's own interval is $[-0.506, +0.357]$, spanning zero. The naive twelve-seed pool, diagnostic only, would read $[-0.509, +0.147]$ and would have classified the transient as refuted; the gate exists because pooling a high-variance three-seed cohort with a lower-variance nine-seed cohort without a variance check assumes exactly what a replication is supposed to test (Figure 3, right, Appendix B).

What stands. The transient is neither confirmed nor refuted; it keeps its original three-seed weight, no more. The evidence leans one way: 7 of 10 primary and 7 of 10 held-out (checkpoint, load) pairs are computable, and every computable pooled interval contains zero. The variance mismatch is not established as a systematic between-wave shift; its direction flips across checkpoints, most consistent with small-n variance imprecision in the archived cohort. A replication wave whose own integrity gate can refuse produces a well-defined result even without a verdict.

6 Related Work

Zoology (Arora et al., 2023) and Based (Arora et al., 2024) establish the recall–state-size tradeoff and the parameter-matched ladder pattern this program's scale series follows; both train architectures for the probe setting, where this program probes general-purpose checkpoints post hoc and reports on instrument validity, not a capacity frontier. Okpekpe & Orvieto (2025) give a positive causal account of recall differences in modern recurrent models; this paper makes no positive mechanism claim and instead bounds one named mechanism three ways. Nichani et al. (2024) supply the associative-memory analysis behind the exact-recovery readout: argmax decoding can succeed without the state holding the claimed structure. The reporting practice follows the NeurIPS pre-registration workshops (Bertinetto et al., 2021) and the negative-results venue tradition (Forde et al., 2020): outcome partitions fixed before unblinding, refusals reported as results.

7 Scope of the Bounds

Bounded. The specific readout construct, one layer's terminal state applied h times to a query under exact continuous recovery, carries no signal in this checkpoint family at any tested scale, surface form, or training regime. Intervention effects larger than 1.5 to 1.7 loss units either way are excluded at the registered confidence, and the transient is bounded in duration, direction, and weight.

Not bounded. The hypothesis survives at effect sizes below the power floor and through mechanisms the readout cannot see, most plausibly cross-layer hand-off, where one layer resolves the first hop and a later layer the next; the readout tests single-layer state self-iteration, a strategy a multi-layer model never trained on the grammar has no particular reason to adopt. Wave 4's dissociation points the same way: the information is present, but not in the form this readout assumes. Nothing here transfers beyond the DeltaNet family or the bind/query grammar.

Lessons for small-scale practice. A registered gate without an enforcing code path is a sentence, not a safeguard; refusal at the process boundary worked twice on real failures. Outcome routing fixed before unblinding kept a wall of zeros from becoming an overclaimed refutation, and a harm-direction transient from becoming an intervention-helps headline. The program cost approximately 9.5 GPU-hours; the scarce commodity was validity, not compute, and the gates purchased it.

References

Simran Arora, Sabri Eyuboglu, Aman Timalsina, Isys Johnson, Michael Poli, James Zou, Atri Rudra, and Christopher Ré. Zoology: Measuring and improving recall in efficient language models. arXiv preprint arXiv:2312.04927, 2023.

Simran Arora, Sabri Eyuboglu, Michael Zhang, Aman Timalsina, Silas Alberti, Dylan Zinsley, James Zou, Atri Rudra, and Christopher Ré. Simple linear attention language models balance the recall-throughput tradeoff. arXiv preprint arXiv:2402.18668, 2024.

Luca Bertinetto, João F. Henriques, Samuel Albanie, Michela Paganini, and Gül Varol. Preface. In NeurIPS 2020 Workshop on Pre-registration in Machine Learning, volume 148 of Proceedings of Machine Learning Research, pp. i–i. PMLR, 2021.

Jessica Zosa Forde, Francisco Ruiz, Melanie F. Pradier, and Aaron Schein (eds.). Proceedings on "I Can't Believe It's Not Better!" at NeurIPS Workshops, volume 137 of Proceedings of Machine Learning Research, 2020. PMLR.

Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models. arXiv preprint arXiv:1609.07843, 2016.

Eshaan Nichani, Jason D. Lee, and Alberto Bietti. Understanding factual recall in transformers via associative memories. arXiv preprint arXiv:2412.06538, 2024.

Destiny Okpekpe and Antonio Orvieto. Revisiting associative recall in modern recurrent models. arXiv preprint arXiv:2508.19029, 2025.

Imanol Schlag, Kazuki Irie, and Jürgen Schmidhuber. Linear transformers are secretly fast weight programmers. arXiv preprint arXiv:2102.11174, 2021.

Songlin Yang, Bailin Wang, Yu Zhang, Yikang Shen, and Yoon Kim. Parallelizing linear transformers with the delta rule over sequence length. arXiv preprint arXiv:2406.06484, 2024.

A Registered Gates and Routing Rules

This appendix reproduces, in compact form, the pre-registered gates and mechanical routing rules the main text relies on. Each was fixed in the program's design registry before the corresponding data existed.

Geometric-readout validity gates (waves 1–4). A cell's readout is valid only if all of the following hold. (a) h=1 sanity floor: the one-hop recovered fraction must exceed the upper edge of a shuffled-pairing null by more than the null's width, and must also exceed an absolute floor of 0.10. (b) Premise (iii): the median cosine between an entity's query projection and the same entity's key projection must exceed the 95th percentile of a shuffled-entity null. (c) Premise (iv): the median cosine between the key and value projections of the same token must exceed the same null percentile. Routing: failure of any gate routes the cell to probe-invalid; the registered rule states that gate failure never routes to refutation.

Behavioral-contrast classification (wave 5). The per-checkpoint arm effect is $\Delta(c) = L_q(\text{control}, c) - L_q(\text{arm}, c)$, with three-seed 95-percent intervals ($t(2, .975) = 4.303$). A checkpoint satisfies the differential condition when the K=32 interval excludes zero, the K=20 interval does not, and $|\Delta_{32}| > |\Delta_{20}|$. Trajectories classify into six registered outcomes: persistent, transient (the condition holds at an interior checkpoint and fails at the terminal one), late-emergent, converged-equivalent, unresolved, and a familiarization-null exclusion. The three-seed power floor (a detectable effect of roughly 1.5 to 1.7 loss units, derived from the control cells' between-seed spread) was registered together with the instruction that unresolved outcomes be reported as measurement bounds.

Replication integrity gates (wave 6). Archived seeds load from stored values only; a runtime guard raises on any live re-scoring of an archived seed. Before pooling, a batch-effect gate requires (a) cohort mean shift below twice the pooled standard error (Welch form) and (b) cohort variance ratio below a pinned 4.0. On gate failure the registered routing reports both cohorts separately, produces no decision-grade pooled interval, and classifies the wave as batch-effect-flagged, outside the four-outcome partition (confirmed at magnitude, confirmed smaller, refuted, new pattern) that a passing gate would have entered.

B Supplementary Figures

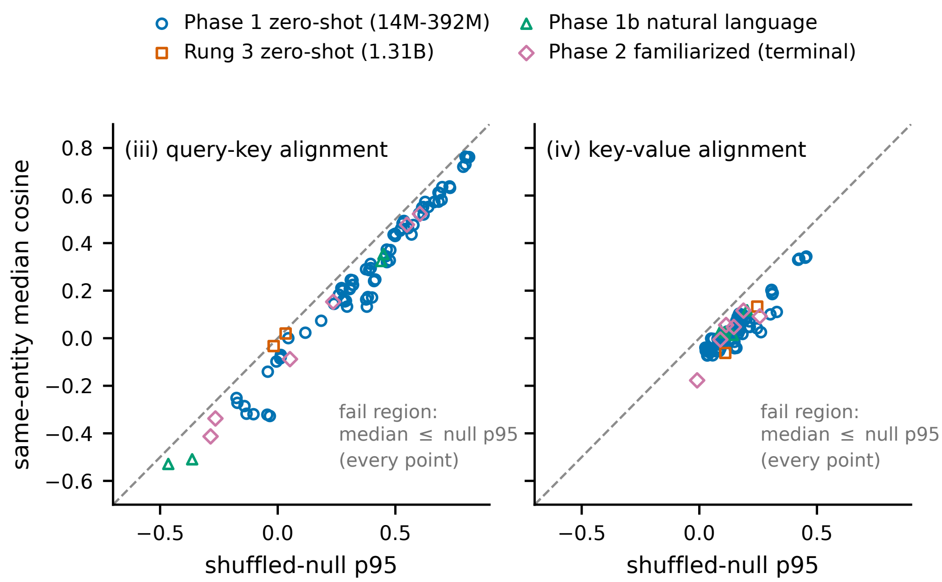

Figure 1: Validity-gate failure at every cell of every instrument. Each point is one cell's same-entity median cosine for premise (iii) (query–key, left) or premise (iv) (key–value, right) against that cell's shuffled-entity null p95; passing requires a median above the dashed identity line. All 90 cells across the four geometric-readout waves fall on or below the line, at every scale from 14M to 1.31B and under every surface form and training regime.

C Reproducibility

Every number in this paper is computed from archived raw result files (per-cell JSON records written by the experiment chains at run time), not from summaries. All figures regenerate from one versioned script that loads only those archived files and asserts the md5 checksum of every file it reads against a recorded manifest, so a changed or substituted artifact fails the build rather than silently altering a panel. The archived raw files, the figure script with its checksum manifest, the experiment chains, and the full pre-registration text will be released with the camera-ready version.

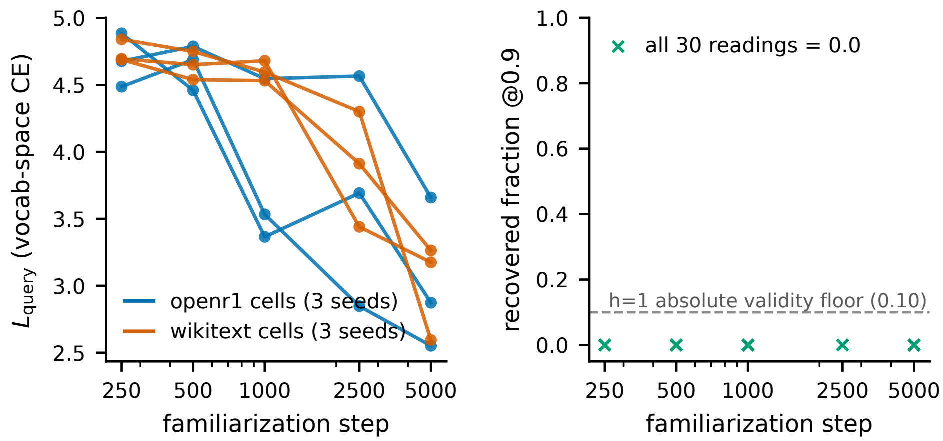

Figure 2: Vocabulary/geometry dissociation under task familiarization (wave 4). Left: held-out in-context query loss L_q for all six control-arm cells (two corpora, three seeds) falls by 21.8 to 46.4 percent over 5,000 familiarization steps. Right: the geometric gate's recovered fraction at the same five checkpoints of the same cells is exactly 0.0 at all 30 readings, never approaching its 0.10 validity floor (dashed). The model learns the task through one readout while the other stays at floor.

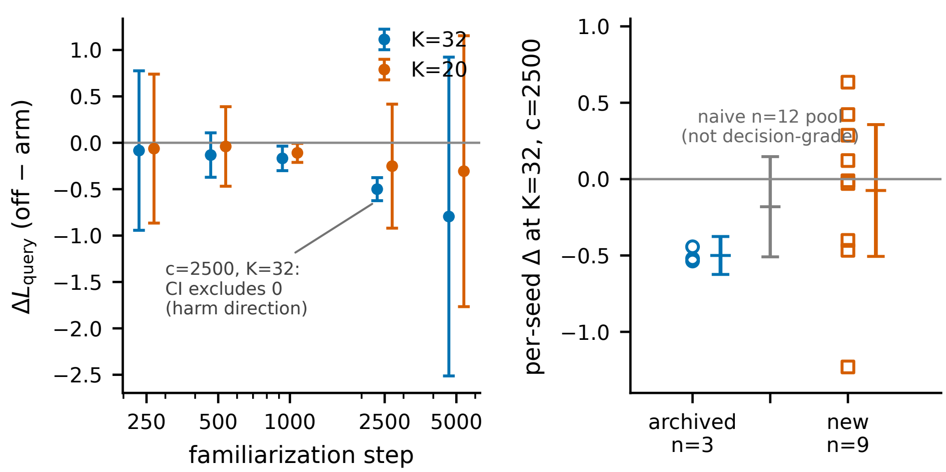

Figure 3: The transient and its stopped replication. Left: per-checkpoint arm effect ΔL_q (control minus per-token arm, encyclopedic corpus) with three-seed 95-percent intervals at both loads; the single reading satisfying the registered differential condition is at checkpoint 2,500, K=32, in the harm direction, and dissolves by checkpoint 5,000. Right: the same anchor cell at n=12; the archived three-seed cohort (circles, interval $[-0.624, -0.376]$) and the new nine-seed cohort (squares, interval $[-0.506, +0.357]$) fail the pre-registered variance-ratio gate (4.47 against a pinned 4.0), so the naive pooled interval (gray, dotted) is diagnostic only and no decision-grade pooled verdict exists.