
Dimension, Not Solvability: Trained Matrix States Converge to the Minimal Faithful Representation Dimension

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 When do differently structured tasks drive a learner to the same representation?
2 We train identical matrix-state sequence models end-to-end on the word prob-
3 lems of five finite groups, S_3 , S_4 , A_5 , S_5 , A_6 , spanning the solvable/non-solvable
4 divide, and measure the effective rank of the learned state against each group’s
5 minimal faithful real representation dimension $d_{\min}$, the exact dimensionality rep-
6 resentation theory assigns the task. The learned rank converges onto the algebraic
7 minimum: group means 1.88/2.85/2.83/3.59/4.74 against $d_{\min} = 2/3/3/4/5$, all
8 19 seeds inside the pre-registered $[0.7, 1.3] \cdot d_{\min}$ band, Spearman $\rho = 0.9747$, the
9 maximum this design permits under the S_4/A_5 tie. The designed head-to-head is
10 that tie: S_4 (solvable) and A_5 (non-solvable) share $d_{\min} = 3$, so if learned dimen-
11 sion were set by computational complexity class the pair should separate, and if
12 set by representation dimension it should coincide. A pre-registered Welch equiv-
13 alence test at ± 0.5 rank-units declares equivalence decisively (difference 0.019,
14 both one-sided tests near seven times the critical value). A pre-registered causal
15 force-rank razor adds a third leg: capping rank one below $d_{\min}$ pins recovery un-
16 der 0.9 by the metric’s own geometry ($\sqrt{(d_{\min}-1)/d_{\min}} < 0.9$ for every group
17 tested), a floor no capped cell violates, and restoring $d_{\min}$ lifts it: recovery empiri-
18 cally returns, past the single-seed anchor in most groups. Learned representational
19 dimension follows the task’s algebra, not its complexity class.

20 1 Introduction

21 A central question for unifying accounts of neural representations is what, exactly, converges: when
22 systems trained on different problems arrive at similar internal structure, which property of the prob-
23 lems predicts the shared solution? We study a setting where the answer can be computed in closed
24 form before training. The word problem of a finite group G , mapping a sequence of generators to
25 its composed product, admits an exact algebraic description of the minimal structure any faithful
26 solution must carry: $d_{\min}(G)$, the dimension of G ’s smallest faithful real linear representation. The
27 loss target already fixes $d_{\min}$ as the informative rank; the open question is whether the model’s own
28 recruited capacity, two ambient dimensions wider than $d_{\min}$, settles there too, tracking the group’s
29 algebra, or instead sorts with the group’s computational class, solvable versus non-solvable, the di-
30 vide at which word problems become NC^1 -complete [Barrington, 1989] and at which state-tracking
31 expressivity results separate recurrent architectures [Merrill et al., 2024, Grazzi et al., 2025].

32 We train one architecture, a Transformer-encoded single-matrix-state model under a hard bottleneck,
33 on five groups spanning that divide, and measure the learned state’s restricted effective rank against
34 $d_{\min}$. Three results. First, convergence to the algebraic minimum: per-group mean rank lands within
35 11% of $d_{\min}$, all 19 seeds inside a pre-registered band, with Spearman $\rho = 0.9747$, the design’s

tie-capped maximum (Section 3). Second, the designed head-to-head at matched dimension: S_4 (solvable) and A_5 (non-solvable) share $d_{\min} = 3$, and a pre-registered equivalence test declares their learned ranks statistically equivalent, so the pair coincides on dimension rather than separating on complexity class (Section 4). Third, the converged dimension is causally load-bearing: capping rank at $d_{\min} - 1$ pins recovery below the readout’s own geometric ceiling in every group, and at $d_{\min}$ that ceiling lifts and recovery empirically returns (Section 5). A measurement lesson runs through all three: two instrument defects each produced a plausible wrong reading before being caught against raw run artifacts (Appendix B), and the convergence law above is what survived the repaired instrument.

2 Task, Model, and Instrument

Task. A word $w = g_{i_1} \cdots g_{i_L}$ is drawn as a random walk on G ’s Cayley graph over a symmetric generating set ($L \sim U\{1, \dots, 8\}$) and must map to $\rho_G(g_{i_1} \cdots g_{i_L})$, the composed product’s image under a pinned orthogonal reference representation of dimension $d_{\min}(G)$. The five groups are S_3, S_4, A_5, S_5, A_6 with $d_{\min} = 2, 3, 3, 4, 5$ (realizations in Appendix A); S_3 and S_4 are solvable, the rest are not.

Model. A Transformer encoder with d_{state} learned row-reader queries writes each word into a single state $Z(w) \in \mathbb{R}^{d_{\text{state}} \times d_{\text{state}}}$, with $d_{\text{state}} = d_{\min} + 2$ per group so that over-recruitment is expressible. The loss is cosine distance to the block-diagonal embedding of the reference ($\rho_G(\cdot) \oplus I_2$ in the observational sweep, $\rho_G(\cdot) \oplus 0$ in the causal wave); the readout is fixed, with no learned weights that could launder rank, and the decoder reads only Z , verified by a gradient blank-out test with a planted-leak negative control. Per-group step budgets (8k–40k) were pinned by pre-registered convergence bars before the decisional sweep ran.

Instrument. The learned state is defined only up to scale and an orthogonal change of basis (a Schur intertwiner): $Z(w) \approx cQ[\rho_G(w) \oplus \cdot]Q^\top$. We estimate the model’s own dominant $d_{\min}$ -subspace U from the SVD of the centered covariance of $Z(w)$ over held-out words, measure **restricted effective rank**, the entropy effective rank of $U^\top Z(w)U$, and score **degauged recovery**, cosine to the reference after fitting (c, Q) on a fitting split and evaluating on a disjoint split, reported as $\text{rec}@0.9$ (the fraction of eval words above cosine 0.9). Rank is invariant under the fitted gauge, so a rank-deficient state cannot be degauged into a full-rank one; centering is load-bearing (Appendix B). The observational band $[0.7, 1.3] \cdot d_{\min}$, the equivalence margin and sample size, the force-rank grid and its recovery bar, and every falsifier were fixed before the decisive runs.

3 Convergence to the Algebraic Minimum

Figure 1 is the headline: unconstrained trained states recruit restricted effective rank 1.877 ± 0.060 , 2.852 ± 0.054 , 2.832 ± 0.062 , 3.591 ± 0.069 , 4.736 ± 0.023 (mean $\pm$ sd over seeds) against $d_{\min} = 2, 3, 3, 4, 5$, within 11% of the algebraic minimum at every group (S_5 the largest deviation, 10.2%), with all 19 seeds inside the pre-registered $[0.7, 1.3] \cdot d_{\min}$ band. Spearman correlation between learned rank and $d_{\min}$ is $\rho = 0.9747$, and this is the maximum the design can express: S_4 and A_5 tie at $d_{\min} = 3$ by construction, capping ρ below 1. Under the exact permutation null, $P(\rho \geq 0.8) = 8/120 \approx 6.7\%$. By pre-registration this observational leg is corroborating: the readout’s $[1, d_{\min}]$ range makes the band’s upper half non-binding by construction (Appendix B), and the dissociation is carried specifically by S_4/A_5 (Section 4), not the five-group trend; the causal weight is carried by Section 5. A disclosed length-robustness split moves per-group mean recovery cosine by at most 0.013 with no group-selective divergence (Appendix B).

4 Equivalence at Matched Dimension

The pair S_4/A_5 was designed into the family as the decisive contrast: the same $d_{\min} = 3$, opposite sides of the solvability divide, the axis on which expressivity theory separates state-tracking architectures [Merrill et al., 2024, Grazi et al., 2025, Siems et al., 2025]. If learned representational dimension were set by complexity class, the pair should separate; if by representation dimension, it should coincide. Because the interesting outcome is a null, the comparison was pre-registered as an equivalence test rather than a difference test: Welch two-one-sided tests on restricted effective




Figure 1: Learned rank converges to the minimal faithful representation dimension. Each point is one seed’s restricted effective rank on the group word problem vs. its group’s $d_{\min}$ (filled: solvable; open: non-solvable), with the pre-registered $[0.7, 1.3] \cdot d_{\min}$ band and the identity line; all 19 seeds are in band. Inset: per-seed $S_4 - A_5$ rank differences at matched $d_{\min} = 3$ against the pre-registered ± 0.5 rank-unit equivalence margin (shaded); the vertical line is the observed mean difference, $+0.019$.

rank at margin ± 0.5 rank-units (half the spacing of the $d_{\min}$ ladder), $n = 5$ seeds per group, with a pre-run power simulation showing 100% power to *reject* equivalence at a true gap of 1.0 rank-units, so a real class effect of one ladder step could not have hidden inside the margin.

The observed difference is $+0.019$ rank-units (se 0.037, df 7.8), and both one-sided tests pass at roughly seven times the critical value ($t = 13.06$ and 14.12 against $t_{\text{crit}} = 1.865$): equivalence is declared decisively. The two groups’ seed clusters are visually coincident in Figure 1 (inset). Convergence here is governed by what the algebra demands, not by how hard the class is to compute.

5 The Causal Check

Convergence and equivalence are observational; the third leg intervenes. The encoder’s output is spectrally truncated to rank $k \in \{d_{\min}-1, d_{\min}, d_{\min}+1\}$ throughout training against the zero-padded target $\rho_G(\cdot) \oplus 0$ of rank exactly $d_{\min}$, alongside an unconstrained anchor from the same family. Pre-registered reading: if $d_{\min}$ is load-bearing, $k = d_{\min}-1$ must fail and $k \geq d_{\min}$ must recover past $0.9 \times$ the anchor’s $\text{rec}@0.9$ on the pre-pinned conservative readout; one below- $d_{\min}$ cell recovering falsifies necessity.

Figure 2 shows a necessity floor at $d_{\min}$, not a plateau. The floor is set by the readout’s own geometry, not discovered by training: $k = d_{\min}-1$ yields $\text{rec}@0.9 = 0.000$ in all five groups and in all four independent seeds of the S_3 extension, exactly as the bound $\sqrt{(d_{\min}-1)/d_{\min}} < 0.9$ requires (Appendix B); the same cells’ recovery cosines (0.61–0.84) sit under that ceiling, evidence the capped models reached their rank-constrained optimum rather than collapsing. Sufficiency is the empirical result: S_4, A_5, S_5, A_6 clear their pre-registered bars (0.800/0.700/0.600/0.650 against 0.585/0.630/0.450/0.585), in a majority of groups against a single noisy anchor seed rather than a stable ceiling, and S_3 , whose decisive cell landed inside a pre-stated ± 0.05 marginality trigger, was extended to four seeds under the trigger’s pre-registered routing: the seed-mean (0.5625) clears the fixed 0.495 bar, though only 2 of 4 seeds individually (Table 2). An earlier 58-cell sweep of this grid nulled through an instrument defect, an ambient identity block in the eye-padded target that taxed every rank-capped arm; the verdict was registered as inconclusive with the tax established from the raw artifacts, and the zero-padded wave above is the registered fix (Appendix B).




Figure 2: The converged dimension is causally load-bearing: per group, crosscheck $\text{rec}@0.9$ on held-out words at train-time force-rank $k \in \{d_{\min}-1, d_{\min}, d_{\min}+1\}$ (one seed per cell), vs. the unconstrained anchor (dashed) and the pre-registered $0.9 \times$ anchor bar (dotted). Recovery is exactly 0.000 at $d_{\min}-1$ in every group, the readout’s geometric floor (Appendix B), and returns at $d_{\min}$. S_3 overlays its four seeds (thin lines; bold mean); its below- $d_{\min}$ zero is unanimous.

6 Related Work and Limitations

Merrill et al. [2024], Grazi et al. [2025], and Siems et al. [2025] frame state tracking by what recurrent architectures *can* express on exactly these word problems, the solvability axis; this work measures what trained states *do* converge to, and finds the learned dimension sorts by representation theory instead. Nazari and Rusch [2026] and Sun et al. [2026] measure state-rank dynamics in pretrained linear-attention models observationally, with no algebraic ground truth to converge to and no causal leg. Nichani et al. [2025] prove associative-memory capacity under argmax decoding, the loophole our exact-continuous-recovery scoring closes; Mishra et al. [2026] train matrix-state RNNs on S_3 composition but measure no state rank and run no rank intervention. Chughtai et al. [2023] reverse-engineer which representation-theoretic structure networks recruit on finite-group composition tasks; this work asks a narrower question on the same substrate, at what dimension, with a pre-registered causal intervention absent there. Huh et al. [2024] propose that models trained on different tasks converge to a shared representation; the present law is a narrower, algebra-driven instance of that convergence.

Limitations. All results are on sub-1M-parameter synthetic models under a deliberate single-state bottleneck; nothing here is a claim about pretrained networks, and the family is five groups with $d_{\min} \in \{2, \dots, 5\}$ at $d_{\text{state}} = d_{\min} + 2$. Causal cells are single-seed except S_3 (four seeds), with the necessity zeros exact wherever seeds exist; two groups at the shortest pinned budget (S_3, S_5) show soft convergence and are also the two with the largest observational deviation and the one marginal causal cell, a pre-registered rather than post hoc budget schedule but a caveat on cross-group comparison strength. The observational readout partially favors rank $\approx d_{\min}$ by construction, which is why it was pre-registered as corroborating and the causal wave carries the claim.

References

- David A. Mix Barrington. Bounded-width polynomial-size branching programs recognize exactly those languages in NC^1 . *Journal of Computer and System Sciences*, 38(1):150–164, 1989.
- Bilal Chughtai, Lawrence Chan, and Neel Nanda. A toy model of universality: Reverse engineering how networks learn group operations. In *International Conference on Machine Learning (ICML)*, 2023. arXiv:2302.03025.
- Riccardo Grazi, Julien Siems, Jörg K. H. Franke, Arber Zela, Frank Hutter, and Massimiliano Pontil. Unlocking state-tracking in linear RNNs through negative eigenvalues. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2411.12537.
- Minyoung Huh, Brian Cheung, Tongzhou Wang, and Phillip Isola. Position: The platonic representation hypothesis. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2405.07987.
- William Merrill, Jackson Petty, and Ashish Sabharwal. The illusion of state in state-space models. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2404.08819.

- 149 Mayank Mishra, Shawn Tan, Ion Stoica, Joseph Gonzalez, and Tri Dao. M²RNN: Non-linear RNNs
 150 with matrix-valued states for scalable language modeling. *arXiv preprint arXiv:2603.14360*,
 151 2026.
- 152 Pedram Nazari and T. Konstantin Rusch. The key to state reduction in linear attention: A rank-based
 153 perspective. *arXiv preprint arXiv:2602.04852*, 2026.
- 154 Eshaan Nichani, Jason D. Lee, and Alberto Bietti. Understanding factual recall in transformers via
 155 associative memories. In *International Conference on Learning Representations (ICLR), Spot-*
 156 *light*, 2025. arXiv:2412.06538.
- 157 Julien Siems, Timur Carstensen, Arber Zela, Frank Hutter, Massimiliano Pontil, and Riccardo
 158 Grazi. DeltaProduct: Improving state-tracking in linear RNNs via householder products. In
 159 *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2502.10297.
- 160 Ang Sun, Hao Zhang, Hang Zhou, Yifei Ma, Yujia Qin, Ting Su, Yang Liu, Zhen Ma, Jian Xu,
 161 Jianfeng Gao, Jiawei Hao, and Ruoxi He. State rank dynamics in linear attention LLMs. *arXiv*
 162 *preprint arXiv:2602.02195*, 2026.

163 A The Five Groups and Reference Representations

G	$ G $	$d_{\min}$	solvable	realization of ρ_G
S_3	6	2	yes	standard (dihedral)
S_4	24	3	yes	cube rotations
A_5	60	3	no	icosahedral rotations
S_5	120	4	no	standard (zero-sum)
A_6	360	5	no	standard (zero-sum)

Table 1: The five task groups. $d_{\min}$ is the minimal faithful real representation dimension; each reference representation is orthogonal by construction and verified faithful (injective on G) at build time. S_4/A_5 form the designed pair: matched $d_{\min}$, opposite solvability.

164 B Instrument Details and the Two Defects

165 **Centering.** For an orthogonal target the uncentered covariance of $Z(w)$ is isotropic and carries no
 166 subspace information; the production lens centers it (a nontrivial irreducible representation has zero
 167 group mean, cancelling the constant block). The step is load-bearing: an uncentered lens scores a
 168 flawless synthetic model 0.705 where the centered lens scores 0.9996 on identical data.

169 **Length-robustness split.** The disclosed split referenced in Section 3 scores only words of length
 170 $L \geq 2$ (the $L = 1$ read is attention-degenerate and was demoted with a mechanism note in the run
 171 records): it moves per-group mean recovery cosine by at most 0.013 (per-seed at most 0.041), with
 172 no group-selective divergence.

173 **The ambient-identity tax.** The first 58-cell sweep’s force-rank target was the eye-padded $\rho_G(\cdot) \oplus I_2$,
 174 whose rank is $d_{\text{state}} = d_{\min} + 2$ for every group. The constant identity block is the cheapest reliable
 175 loss reduction, so a rank-capped arm buys it first, leaving effective budget $k - 2 < d_{\min}$ for ρ_G at
 176 every grid point; no capped cell in that sweep ever tested the confirm direction. The raw-artifact
 177 signature: the 39 force-rank cells’ direct cosine matched the rank- k optimum $\sqrt{k/d_{\text{state}}}$ with mean
 178 absolute deviation 0.028 (maximum 0.166; the two S_5 below- $d_{\min}$ cells are the only outliers, a
 179 disclosed optimization shortfall), so the arms had trained to their rank-constrained ceiling rather
 180 than under-training. The sweep verdict was registered as inconclusive with the tax as mechanism;
 181 the zero-padded causal wave of Section 5 is the registered fix, its 30 cells configuration-verified
 182 one-by-one against a manifest re-derived independently from the design record. In that wave an
 183 eye-padded corroboration arm reproduces the failure on demand at raw $k = d_{\min} + 1$ (effective
 184 budget $d_{\min} - 1$) while the tax-paid point recovers.

185 **The observational band’s non-binding upper half.** `entity_subspace_from_words` returns the
 186 top- $d_{\min}$ left singular vectors of the centered covariance by construction, so the restricted state

187 $U^\top ZU$ is always $d_{\min} \times d_{\min}$, and the entropy effective rank of a $d_{\min} \times d_{\min}$ input has range
 188 $[1, d_{\min}]$. Restricted effective rank therefore cannot exceed $d_{\min}$, and so cannot exceed $1.3 \cdot d_{\min}$,
 189 for any seed of any group; only the band’s lower half, $\geq 0.7 \cdot d_{\min}$, is a test the data could fail. Every
 190 group mean lands in $[0.895, 0.951] \cdot d_{\min}$.

191 **The rank-constrained cosine ceiling.** The causal wave’s target is $\rho_G(\cdot) \oplus 0$, of rank exactly
 192 $d_{\min}$ with all $d_{\min}$ informative singular values equal to 1 (the reference representation is orthog-
 193 onal by construction, Appendix A). For any matrix M of rank $\leq k \leq d_{\min}$, von Neumann’s
 194 trace inequality gives $\langle M, T \rangle_F \leq \sum_{i=1}^k \sigma_i(M) \sigma_i(T) = \sum_{i=1}^k \sigma_i(M)$, and Cauchy–Schwarz
 195 bounds $\sum_{i=1}^k \sigma_i(M) \leq \sqrt{k} \|M\|_F$, so $\cos(M, T) \leq \sqrt{k/d_{\min}}$. At $k = d_{\min} - 1$ this ceiling is
 196 $0.707/0.816/0.816/0.866/0.894$ for $S_3/S_4/A_5/S_5/A_6$, below the rec@0.9 threshold of 0.9 in ev-
 197 ery case: no rank- $(d_{\min} - 1)$ state, however trained, can register a single word above cosine 0.9
 198 against this target. The necessity result ($\text{rec@0.9} = 0.000$ at $d_{\min} - 1$, all groups, all seeds) is this
 199 bound realized, not an independent discovery about what SGD finds; the corresponding empirical
 200 content is that the observed cosines at $d_{\min} - 1$ (0.61–0.84) sit under their respective rank-constrained
 201 optima, evidence optimization reaches the geometry’s own ceiling. At $k \geq d_{\min}$ the ceiling is 1.0
 202 (exact recovery is representable), so recovery above 0 at $d_{\min}$ is not guaranteed by the metric and is
 203 the paper’s genuine causal finding (sufficiency).

204 **Primary and crosscheck degauging.** The design record’s default pipeline treats scale-only degaug-
 205 ing ($\hat{Q} = I$) as the primary metric, with the fitted- (c, Q) Procrustes score retained as a robustness
 206 cross-check, after calibration on unconstrained checkpoints found $\hat{Q} \approx I$ empirically. Under the
 207 force-rank grid’s zero-padded target this equivalence breaks: the informative block’s degenerate
 208 singular spectrum makes the scale-only fit’s basis arbitrary, so the pre-registration for the causal
 209 wave specifically pins the fitted- (c, Q) score, denoted rec@0.9 throughout Section 5 and Figure 2,
 210 as decisional; no conclusion in this paper reads the scale-only score. The two diverge sharply on this
 211 grid, which is why the distinction is disclosed here.

212 C S_3 Per-Seed Causal Detail

seed	anchor	$k=d_{\min}-1$	$k=d_{\min}$	$k=d_{\min}+1$	own bar	clears
0	0.550	0.000	0.450	0.550	0.495	no (−0.045)
1	0.600	0.000	0.550	0.750	0.540	yes (+0.010)
2	0.800	0.000	0.600	0.750	0.720	no (−0.120)
3	0.600	0.000	0.650	0.600	0.540	yes (+0.110)

Table 2: S_3 four-seed causal detail (crosscheck rec@0.9) behind the seed-mean confirmation in Section 5; the own bar is $0.9 \times$ each seed’s anchor. $k = d_{\min} - 1$ reads exactly 0.000 in all four seeds; $k = d_{\min}$ clears each seed’s own bar in 2 of 4 (seeds 1, 3), missing in the other 2 because the anchor itself ranges 0.550–0.800 across seeds. The pre-registered criterion compares the seed-mean $k = d_{\min}$ value (0.5625) against the bar fixed before the extension ran (0.495, from seed 0’s anchor), specifically to avoid computing the bar from the same noisy seeds it is judged against; recomputing the bar from the four-seed anchor mean (0.574) would put the seed-mean value 0.011 below it.

G	$d_{\min}$	n	restricted eff. rank	band	in band
S_3	2	3	1.877 ± 0.060	$[1.4, 2.6]$	3/3
S_4	3	5	2.852 ± 0.054	$[2.1, 3.9]$	5/5
A_5	3	5	2.832 ± 0.062	$[2.1, 3.9]$	5/5
S_5	4	3	3.591 ± 0.069	$[2.8, 5.2]$	3/3
A_6	5	3	4.736 ± 0.023	$[3.5, 6.5]$	3/3

Table 3: Restricted effective rank of the unconstrained trained state (mean $\pm$ sd over seeds) against each group’s $d_{\min}$, with the pre-registered $[0.7, 1.3] \cdot d_{\min}$ band. Every seed of every group is in band; Spearman $\rho = 0.9747$ is the tie-capped maximum; the dimension-matched pair S_4/A_5 differs by 0.019 rank-units.